Problem Statement

Ola relies heavily on its driver network to provide efficient ride-sharing services. However, driver attribution (churn) poses a significant challenge as it affects service availability, operatonal costs, and overall revenue. The objective of this case study it analyze historical driver data to identify patterns associated with driver churn and build a predictive model using ensemble learning techniques. The insights from this analysis will help Ola proactively identify drivers at risk of leaving and implement strategies to improve driver retention.

Objective:

- Analyze driver data to understand patterns related to driver attrition.
- Identify key factors that influence driver churn such as income, ratings, and business value.
- Build predictive models using machine learning and ensemble techniques to identify drivers likely to churn.
- Evaluate model performance using appropriate metrics to ensure reliable predictions.
- Generate actionable insights to help the company reduce driver attrition.
- Support business decision-making by recommending strategies to improve driver retention.

Dataset Overview:

- MMMM-YY: Reporting month and year.
- Driver_ID: A unique identifier for every driver.
- Age: Age of the driver.
- Gender: Driver's gender. Male: 0, Female: 1.
- City: City code representing the city the driver operates in.
- Education_Level: Education level of the driver, categorized into 0 for 10+, 1 for 12+, and 2 for graduate.
- Income: Average monthly income of the driver.
- Date Of Joining: The date when the driver joined Ola.
- LastWorkingDate: The most recent or final day the driver worked with Ola.
- Joining Designation: Designation of the driver at the onset of their journey with Ola.
- Grade: A grade assigned to the driver at the reporting time, likely denoting performance or other metrics.
- Total Business Value: The total monetary value (business) a driver brings in a month. Negative values might indicate cancellations, refunds, or other financial adjustments.
- Quarterly Rating: Rating assigned to drivers on a quarterly basis. Ratings range from 1 to 5, with 5 being the best.

```
!wget "https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/002/492/original/ola_driver_scaler.csv" -O ola_driver.csv
```

```
--2026-03-27 18:55:37--
https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/002/492
```

```
/original/ola_driver_scaler.csv
Resolving d2beiqkhq929f0.cloudfront.net
(d2beiqkhq929f0.cloudfront.net)... 13.225.41.18, 13.225.41.124,
13.225.41.45, ...
Connecting to d2beiqkhq929f0.cloudfront.net
(d2beiqkhq929f0.cloudfront.net)|13.225.41.18|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1127673 (1.1M) [text/plain]
Saving to: ‘ola_driver.csv’

 ola_driver.csv        0%[                    ]       0  --.-KB/s
ola_driver.csv      100%[===================>]   1.08M  --.-KB/s    in
0.05s

2026-03-27 18:55:38 (20.1 MB/s) - ‘ola_driver.csv’ saved
[1127673/1127673]
```

```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('ola_driver.csv')
df.head()
```

```
/usr/local/lib/python3.12/dist-packages/google/colab/
_dataframe_summarizer.py:88: UserWarning: Could not infer format, so
each element will be parsed individually, falling back to `dateutil`.
To ensure parsing is consistent and as-expected, please specify a
format.
  cast_date_col = pd.to_datetime(column, errors="coerce")
```


```
{"summary":"{\n  \"name\": \"df\",\n  \"rows\": 19104,\n  \"fields\":
[\n    {\n      \"column\": \"Unnamed: 0\",\n      \"properties\": {\n
\"dtype\": \"number\",\n        \"std\": 5514,\n        \"min\": 0,\n
\"max\": 19103,\n        \"num_unique_values\": 19104,\n
\"samples\": [\n          18299,\n          9376,\n          4518\n
],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n
}\n    },\n    {\n      \"column\": \"MMM-YY\",\n      \"properties\":
{\n        \"dtype\": \"object\",\n        \"num_unique_values\": 24,\
n        \"samples\": [\n          \"03/01/20\",\n
\"10/01/19\",\n          \"01/01/19\"\n        ],\n
\"semantic_type\": \"\",\n        \"description\": \"\"\n      }\
n    },\n    {\n      \"column\": \"Driver_ID\",\n
\"properties\": {\n        \"dtype\": \"number\",\n        \"std\":
810,\n        \"min\": 1,\n        \"max\": 2788,\n
\"num_unique_values\": 2381,\n        \"samples\": [\n          1663,\
n          1264,\n          1618\n        ],\n
\"semantic_type\": \"\",\n        \"description\": \"\"\n      }\
n    },\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n
```

\"dtype\": \"number\",\n   \"std\": 6.2579116861907345,\n
\"min\": 21.0,\n   \"max\": 58.0,\n   \"num_unique_values\":
36,\n   \"samples\": [\n   58.0,\n   41.0,\n
24.0\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"Gender\",\n   \"properties\": {\n   \"dtype\": \"number\",\n
\"std\": 0.4933670037660394,\n   \"min\": 0.0,\n   \"max\":
1.0,\n   \"num_unique_values\": 2,\n   \"samples\": [\n
1.0,\n   0.0\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"City\",\n   \"properties\": {\n   \"dtype\": \"category\",\n
\"num_unique_values\": 29,\n   \"samples\": [\n
\"C22\",\n   \"C5\"\n   ],\n   \"semantic_type\":
\"\",\n   \"description\": \"\"\n   }\n   },\n   {\n
\"column\": \"Education_Level\",\n   \"properties\": {\n
\"dtype\": \"number\",\n   \"std\": 0,\n   \"min\": 0,\n
\"max\": 2,\n   \"num_unique_values\": 3,\n   \"samples\":
[\n   2,\n   0\n   ],\n   \"semantic_type\":
\"\",\n   \"description\": \"\"\n   }\n   },\n   {\n
\"column\": \"Income\",\n   \"properties\": {\n   \"dtype\":
\"number\",\n   \"std\": 30914,\n   \"min\": 10747,\n
\"max\": 188418,\n   \"num_unique_values\": 2383,\n
\"samples\": [\n   44273,\n   35370\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"Dateofjoining\",\n
\"properties\": {\n   \"dtype\": \"object\",\n
\"num_unique_values\": 869,\n   \"samples\": [\n
\"14/09/19\",\n   \"01/06/18\"\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"LastWorkingDate\",\n
\"properties\": {\n   \"dtype\": \"date\",\n   \"min\":
\"2018-12-31 00:00:00\",\n   \"max\": \"2020-12-28 00:00:00\",\n
\"num_unique_values\": 493,\n   \"samples\": [\n
\"05/03/20\",\n   \"10/01/19\"\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"Joining Designation\",\n
\"properties\": {\n   \"dtype\": \"number\",\n   \"std\":
0,\n   \"min\": 1,\n   \"max\": 5,\n
\"num_unique_values\": 5,\n   \"samples\": [\n   2,\n
5\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"Grade\",\n   \"properties\": {\n   \"dtype\": \"number\",\n
\"std\": 1,\n   \"min\": 1,\n   \"max\": 5,\n
\"num_unique_values\": 5,\n   \"samples\": [\n   2,\n
5\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"Total Business Value\",\n   \"properties\": {\n   \"dtype\":
\"number\",\n   \"std\": 1128312,\n   \"min\": -6000000,\n
\"max\": 33747720,\n   \"num_unique_values\": 10181,\n

\"samples\": [\n 431090,\n 720180\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n {\n \"column\": \"Quarterly Rating\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 1,\n \"min\": 1,\n \"max\": 4,\n \"num_unique_values\": 4,\n \"samples\": [\n 1,\n 3\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n]\n}","type":"dataframe","variable_name":"df"}

Data Understanding

```
df.shape
```

```
(19104, 14)
```

```
df.describe()
```

{"summary":"{\n \"name\": \"df\",\n \"rows\": 8,\n \"fields\": [\n {\n \"column\": \"Unnamed: 0\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6884.95505562544,\n \"min\": 0.0,\n \"max\": 19104.0,\n \"num_unique_values\": 7,\n \"samples\": [\n 19104.0,\n 9551.5,\n 14327.25\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n {\n \"column\": \"Driver_ID\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6344.561458172487,\n \"min\": 1.0,\n \"max\": 19104.0,\n \"num_unique_values\": 8,\n \"samples\": [\n 1415.5911327470687,\n 1417.0,\n 19104.0\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n {\n \"column\": \"Age\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6721.473927583281,\n \"min\": 6.2579116861907345,\n \"max\": 19043.0,\n \"num_unique_values\": 8,\n \"samples\": [\n 34.668434595389385,\n 34.0,\n 19043.0\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n {\n \"column\": \"Gender\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6735.752126331659,\n \"min\": 0.0,\n \"max\": 19052.0,\n \"num_unique_values\": 5,\n \"samples\": [\n 0.4187486878018056,\n 1.0,\n 0.4933670037660394\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n {\n \"column\": \"Education_Level\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6753.939461727606,\n \"min\": 0.0,\n \"max\": 19104.0,\n \"num_unique_values\": 6,\n \"samples\": [\n 19104.0,\n 1.0216708542713568,\n 2.0\n],\n

\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\n   },\n   {\n   \"column\": \"Income\",\n   \"properties\": {\n   \"dtype\": \"number\",\n   \"std\": 56436.738053256035,\n   \"min\": 10747.0,\n   \"max\": 188418.0,\n   \"num_unique_values\": 8,\n   \"samples\": [\n 65652.02512562813,\n   60087.0,\n   19104.0\n   ],\n   \"semantic_type\": \"\",\n   \"description\": \"\"\n }\n   },\n   {\n   \"column\": \"Joining Designation\",\n \"properties\": {\n   \"dtype\": \"number\",\n   \"std\": 6753.651377555244,\n   \"min\": 0.8369837171189233,\n \"max\": 19104.0,\n   \"num_unique_values\": 6,\n \"samples\": [\n   19104.0,\n   1.690536013400335,\n 5.0\n   ],\n   \"semantic_type\": \"\",\n \"description\": \"\"\n   }\n   },\n   {\n   \"column\": \"Grade\",\n   \"properties\": {\n   \"dtype\": \"number\",\n \"std\": 6753.512393804842,\n   \"min\": 1.0,\n   \"max\": 19104.0,\n   \"num_unique_values\": 7,\n   \"samples\": [\n 19104.0,\n   2.2526695979899496,\n   3.0\n   ],\n \"semantic_type\": \"\",\n   \"description\": \"\"\n   }\n   },\n   {\n   \"column\": \"Total Business Value\",\n \"properties\": {\n   \"dtype\": \"number\",\n   \"std\": 12313845.373558126,\n   \"min\": -6000000.0,\n   \"max\": 33747720.0,\n   \"num_unique_values\": 8,\n   \"samples\": [\n   571662.074958124,\n   250000.0,\n 19104.0\n   ],\n   \"semantic_type\": \"\",\n \"description\": \"\"\n   }\n   },\n   {\n   \"column\": \"Quarterly Rating\",\n   \"properties\": {\n   \"dtype\": \"number\",\n   \"std\": 6753.57600547721,\n   \"min\": 1.0,\n   \"max\": 19104.0,\n   \"num_unique_values\": 7,\n \"samples\": [\n   19104.0,\n   2.008898659966499,\n 3.0\n   ],\n   \"semantic_type\": \"\",\n \"description\": \"\"\n   }\n   }\n]\n}","type":"dataframe"}

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  ------                --------------  -----
 0   Unnamed: 0            19104 non-null  int64
 1   MMM-YY                19104 non-null  object
 2   Driver_ID             19104 non-null  int64
 3   Age                   19043 non-null  float64
 4   Gender                19052 non-null  float64
 5   City                  19104 non-null  object
 6   Education_Level       19104 non-null  int64
 7   Income                19104 non-null  int64
 8   Dateofjoining         19104 non-null  object
 9   LastWorkingDate       1616 non-null   object
```

```
 10  Joining Designation   19104 non-null  int64
 11  Grade                 19104 non-null  int64
 12  Total Business Value  19104 non-null  int64
 13  Quarterly Rating      19104 non-null  int64
dtypes: float64(2), int64(8), object(4)
memory usage: 2.0+ MB
```

```
df.isnull().sum()
```

```
Unnamed: 0                  0
MMM-YY                      0
Driver_ID                   0
Age                        61
Gender                     52
City                        0
Education_Level             0
Income                      0
Dateofjoining               0
LastWorkingDate         17488
Joining Designation         0
Grade                       0
Total Business Value        0
Quarterly Rating            0
dtype: int64
```

```
df.dtypes
```

```
Unnamed: 0                 int64
MMM-YY                    object
Driver_ID                  int64
Age                      float64
Gender                   float64
City                      object
Education_Level            int64
Income                     int64
Dateofjoining             object
LastWorkingDate           object
Joining Designation        int64
Grade                      int64
Total Business Value       int64
Quarterly Rating           int64
dtype: object
```

```
df['Dateofjoining'] =
pd.to_datetime(df['Dateofjoining'],format='mixed')
df['LastWorkingDate'] =
pd.to_datetime(df['LastWorkingDate'],format='mixed')
df['MMM-YY'] = pd.to_datetime(df['MMM-YY'],dayfirst=True)
```

```
/tmp/ipykernel_18366/2611634496.py:3: UserWarning: Could not infer
format, so each element will be parsed individually, falling back to
```

```
`dateutil`. To ensure parsing is consistent and as-expected, please
specify a format.
  df['MMM-YY'] = pd.to_datetime(df['MMM-YY'],dayfirst=True)
```

- The dataset contains driver-level monthly records including demographic, performance, and income-related attributes.
- Key features include Age, City, Income, Total Business Value, and Quarterly Rating.
- The dataset does not contain an explicit target variable for churn, which needs to be derived.
- Missing values are observed in columns like Age, gendr, LastWorkingDate.
- Data types include both numerical and categorical variables.
- Since the dataset does not have a predefined target variable, we will create a churn variable using the LastWorkingDate column.
- Drivers with a LastWorkingDate are considered churned (1), otherwise active (0).

```
df = df.sort_values(by=['Driver_ID', 'MMM-YY'])

income_change = df.groupby('Driver_ID')['Income'].last() - \
                df.groupby('Driver_ID')['Income'].first()

rating_change = df.groupby('Driver_ID')['Quarterly Rating'].last() - \
                df.groupby('Driver_ID')['Quarterly Rating'].first()

avg_income = df.groupby('Driver_ID')['Income'].mean()

avg_rating = df.groupby('Driver_ID')['Quarterly Rating'].mean()

income_std = df.groupby('Driver_ID')['Income'].std().fillna(0)

business_value_std = df.groupby('Driver_ID')['Total Business
Value'].std().fillna(0)

active_months = df.groupby('Driver_ID')['MMM-YY'].nunique()
```

Aggregate data by Driver_ID

```
driver_df = df.groupby('Driver_ID').agg({
    'MMM-YY': 'first',
    'Age': 'first',
    'Gender': 'first',
    'City': 'first',
    'Education_Level': 'first',
    'Income': 'mean',
    'Dateofjoining': 'first',
    'LastWorkingDate': 'last',
    'Joining Designation': 'first',
    'Grade': 'last',
    'Total Business Value': 'sum',
```

```
    'Quarterly Rating': 'mean'
}).reset_index()

driver_df = driver_df.merge(income_change.rename('income_change'),
on='Driver_ID')
driver_df = driver_df.merge(rating_change.rename('rating_change'),
on='Driver_ID')
driver_df = driver_df.merge(avg_income.rename('avg_income'),
on='Driver_ID')
driver_df = driver_df.merge(avg_rating.rename('avg_rating'),
on='Driver_ID')
driver_df = driver_df.merge(income_std.rename('income_std'),
on='Driver_ID')
driver_df =
driver_df.merge(business_value_std.rename('business_value_std'),
on='Driver_ID')
driver_df = driver_df.merge(active_months.rename('active_months'),
on='Driver_ID')

driver_df['income_std'] = driver_df['income_std'].fillna(0)
driver_df['business_value_std'] =
driver_df['business_value_std'].fillna(0)

driver_df.head()
```


{"summary":"{\n \"name\": \"driver_df\",\n \"rows\": 2381,\n
\"fields\": [\n {\n \"column\": \"Driver_ID\",\n
\"properties\": {\n \"dtype\": \"number\",\n \"std\":
806,\n \"min\": 1,\n \"max\": 2788,\n
\"num_unique_values\": 2381,\n \"samples\": [\n 1663,\
n 1264,\n 1618\n],\n
\"semantic_type\": \"\",\n \"description\": \"\"\n }\
n },\n {\n \"column\": \"MMM-YY\",\n \"properties\":
{\n \"dtype\": \"date\",\n \"min\": \"2019-01-01
00:00:00\",\n \"max\": \"2020-01-12 00:00:00\",\n
\"num_unique_values\": 24,\n \"samples\": [\n \"2019-
01-10 00:00:00\",\n \"2020-01-04 00:00:00\",\n
\"2019-01-01 00:00:00\"\n],\n \"semantic_type\": \"\",\
n \"description\": \"\"\n }\n },\n {\n
\"column\": \"Age\",\n \"properties\": {\n \"dtype\":
\"number\",\n \"std\": 5.840686217594289,\n \"min\":
21.0,\n \"max\": 58.0,\n \"num_unique_values\": 35,\n
\"samples\": [\n 51.0,\n 41.0,\n 24.0\n
],\n \"semantic_type\": \"\",\n \"description\": \"\"\n
}\n },\n {\n \"column\": \"Gender\",\n \"properties\":
{\n \"dtype\": \"number\",\n \"std\":
0.4919972321279086,\n \"min\": 0.0,\n \"max\": 1.0,\n
\"num_unique_values\": 2,\n \"samples\": [\n 1.0,\n
0.0\n],\n \"semantic_type\": \"\",\n
\"description\": \"\"\n }\n },\n {\n \"column\":

\"City\",\n   \"properties\": {\n   \"dtype\": \"category\",\n
\"num_unique_values\": 29,\n   \"samples\": [\n
\"C22\",\n   \"C5\"\n   ],\n   \"semantic_type\":
\"\",\n   \"description\": \"\"\n   }\n   },\n   {\n
\"column\": \"Education_Level\",\n   \"properties\": {\n
\"dtype\": \"number\",\n   \"std\": 0,\n   \"min\": 0,\n
\"max\": 2,\n   \"num_unique_values\": 3,\n   \"samples\":
[\n   2,\n   0\n   ],\n   \"semantic_type\":
\"\",\n   \"description\": \"\"\n   }\n   },\n   {\n
\"column\": \"Income\",\n   \"properties\": {\n   \"dtype\":
\"number\",\n   \"std\": 28298.214011922697,\n   \"min\":
10747.0,\n   \"max\": 188418.0,\n   \"num_unique_values\":
2339,\n   \"samples\": [\n   58943.0,\n   96076.0\n
],\n   \"semantic_type\": \"\",\n   \"description\": \"\"\n
}\n   },\n   {\n   \"column\": \"Dateofjoining\",\n
\"properties\": {\n   \"dtype\": \"date\",\n   \"min\":
\"2013-04-01 00:00:00\",\n   \"max\": \"2020-12-28 00:00:00\",\n
\"num_unique_values\": 869,\n   \"samples\": [\n   \"2019-
09-14 00:00:00\",\n   \"2018-01-06 00:00:00\"\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"LastWorkingDate\",\n
\"properties\": {\n   \"dtype\": \"date\",\n   \"min\":
\"2018-12-31 00:00:00\",\n   \"max\": \"2020-12-28 00:00:00\",\n
\"num_unique_values\": 493,\n   \"samples\": [\n   \"2020-
05-03 00:00:00\",\n   \"2019-10-01 00:00:00\"\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"Joining Designation\",\n
\"properties\": {\n   \"dtype\": \"number\",\n   \"std\":
0,\n   \"min\": 1,\n   \"max\": 5,\n
\"num_unique_values\": 5,\n   \"samples\": [\n   2,\n
5\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"Grade\",\n   \"properties\": {\n   \"dtype\": \"number\",\n
\"std\": 0,\n   \"min\": 1,\n   \"max\": 5,\n
\"num_unique_values\": 5,\n   \"samples\": [\n   2,\n
5\n   ],\n   \"semantic_type\": \"\",\n
\"description\": \"\"\n   }\n   },\n   {\n   \"column\":
\"Total Business Value\",\n   \"properties\": {\n   \"dtype\":
\"number\",\n   \"std\": 9127115,\n   \"min\": -1385530,\n
\"max\": 95331060,\n   \"num_unique_values\": 1629,\n
\"samples\": [\n   479380,\n   1628780\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"Quarterly Rating\",\n
\"properties\": {\n   \"dtype\": \"number\",\n   \"std\":
0.7196519783986751,\n   \"min\": 1.0,\n   \"max\": 4.0,\n
\"num_unique_values\": 163,\n   \"samples\": [\n
3.1666666666666665,\n   3.625\n   ],\n
\"semantic_type\": \"\",\n   \"description\": \"\"\n   }\
n   },\n   {\n   \"column\": \"income_change\",\n

\"properties\": {\n \"dtype\": \"number\",\n \"std\": 968,\n \"min\": 0,\n \"max\": 12155,\n \"num_unique_values\": 44,\n \"samples\": [\n 8744,\n 6888\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"rating_change\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 0,\n \"min\": -3,\n \"max\": 3,\n \"num_unique_values\": 7,\n \"samples\": [\n 0,\n 1\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"avg_income\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 28298.214011922697,\n \"min\": 10747.0,\n \"max\": 188418.0,\n \"num_unique_values\": 2339,\n \"samples\": [\n 58943.0,\n 96076.0\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"avg_rating\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 0.7196519783986751,\n \"min\": 1.0,\n \"max\": 4.0,\n \"num_unique_values\": 163,\n \"samples\": [\n 3.1666666666666665,\n 3.625\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"income_std\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 380.6438169817499,\n \"min\": 0.0,\n \"max\": 5563.884994487583,\n \"num_unique_values\": 45,\n \"samples\": [\n 2273.11263804005,\n 3046.740581984794\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"business_value_std\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 589340.756286116,\n \"min\": 0.0,\n \"max\": 8074168.089866569,\n \"num_unique_values\": 1656,\n \"samples\": [\n 173233.2992886113,\n 378122.47291056375\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n },\n {\n \"column\": \"active_months\",\n \"properties\": {\n \"dtype\": \"number\",\n \"std\": 6,\n \"min\": 1,\n \"max\": 24,\n \"num_unique_values\": 24,\n \"samples\": [\n 16,\n 15\n],\n \"semantic_type\": \"\",\n \"description\": \"\"\n }\n }\n]\n}","type":"dataframe","variable_name":"driver_df"}

```
driver_df.shape
```

```
(2381, 20)
```

Since the dataset contains monthly records for each driver, the same driver appears multiple times. To analyze driver-level behavior and predict churn effectively, the data was aggregated at the Driver_ID level using appropriate aggregation techniques such as mean, sum, first, and last for different features.

```
df['MMM-YY'] = pd.to_datetime(df['MMM-YY'])

latest_date = df['MMM-YY'].max()

driver_last_date = df.groupby('Driver_ID')['MMM-
YY'].max().reset_index()
driver_last_date.rename(columns={'MMM-YY': 'last_active_date'},
inplace=True)

for days in [60, 90, 120]:
    driver_last_date[f'churn_{days}'] = (
        (latest_date - driver_last_date['last_active_date']).dt.days >
days
    ).astype(int)

for col in ['churn_60', 'churn_90', 'churn_120']:
    print(f"\n{col} distribution:")
    print(driver_last_date[col].value_counts(normalize=True))
```

```
churn_60 distribution:
churn_60
0    0.651407
1    0.348593
Name: proportion, dtype: float64

churn_90 distribution:
churn_90
0    0.651407
1    0.348593
Name: proportion, dtype: float64

churn_120 distribution:
churn_120
0    0.651407
1    0.348593
Name: proportion, dtype: float64
```

```
driver_last_date['churn'] = driver_last_date['churn_90']

driver_df = driver_df.merge(
    driver_last_date[['Driver_ID', 'churn']],
    on='Driver_ID',
    how='left'
)

print(driver_df[['Driver_ID', 'churn']].head())
```

```
   Driver_ID  churn
0          1      1
1          2      0
2          4      0
```

```
3           5       1
4           6       0
```

```
print(driver_df['churn'].value_counts())
```

```
churn
0    1551
1     830
Name: count, dtype: int64
```

```
driver_df = driver_df.dropna(subset=['churn'])
```

- Churn was defined based on driver inactivity using multiple time windows (60, 90, and 120 days).
- All definitions resulted in identical churn distribution, indicating clear separation between active and inactive drivers.
- A 90-day inactivity threshold was selected as the final churn definition for consistency and business interpretability.

```
driver_df.groupby('MMM-YY')['churn'].mean().plot()
```

```
<Axes: xlabel='MMM-YY'>
```

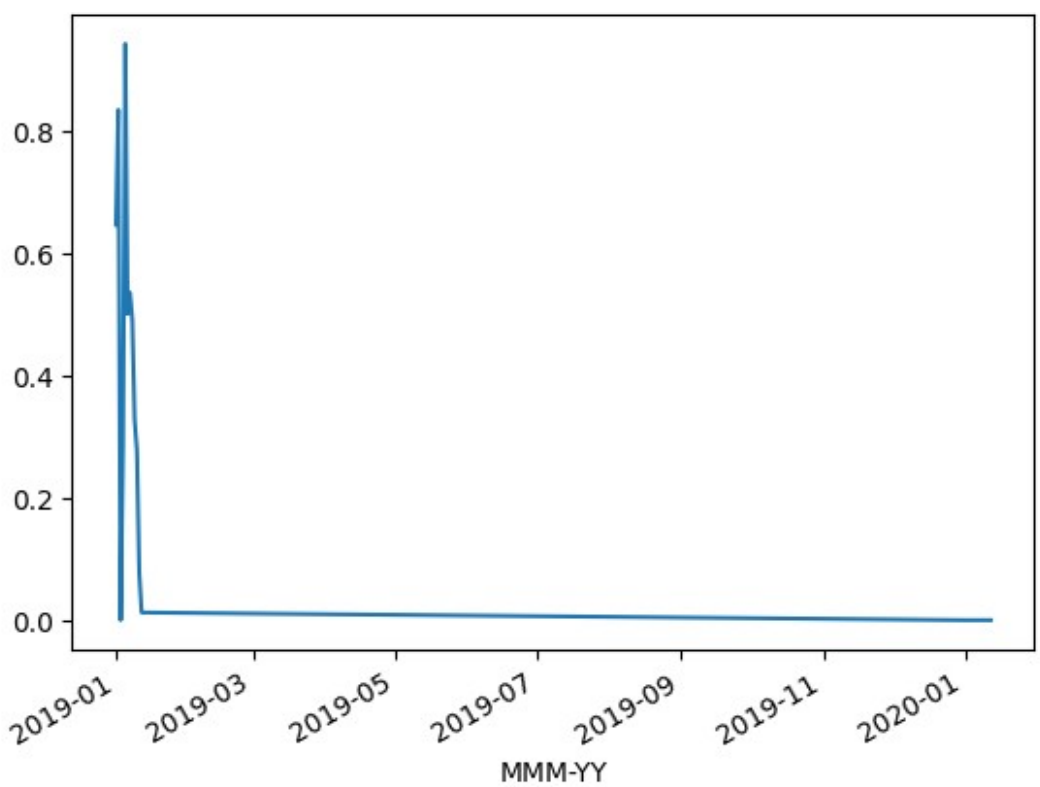


```
driver_df.groupby('MMM-YY')['churn'].value_counts()
```

```
MMM-YY      churn
2019-01-01  1        660
            0        362
2019-01-02  1          5
            0          1
2019-01-03  0          1
2019-01-04  0          5
            1          2
2019-01-05  1         16
            0          1
2019-01-06  0         29
            1         29
2019-01-07  1         46
            0         40
2019-01-08  0         28
            1         27
2019-01-09  0         37
            1         18
2019-01-10  0         47
            1         18
2019-01-11  0         96
            1          8
2019-01-12  0         79
            1          1
2020-01-01  0         56
2020-01-02  0         49
2020-01-03  0         36
2020-01-04  0         71
2020-01-05  0         83
2020-01-06  0         73
2020-01-07  0        100
2020-01-08  0         94
2020-01-09  0         66
2020-01-10  0         70
2020-01-11  0         49
2020-01-12  0         78
Name: count, dtype: int64
```

- Temporal analysis revealed that churn events were concentrated in early months, with later periods containing only active drivers.
- To ensure meaningful model evaluation, the time-based split was adjusted so that both churn and non-churn cases are present in the test set.
- This highlights the importance of understanding temporal data distribution before model validation.

#Univariate & Bivariate Analysis

```
sns.histplot(driver_df['Age'],kde=True)
plt.title('Distribution of Age')
plt.show()
```

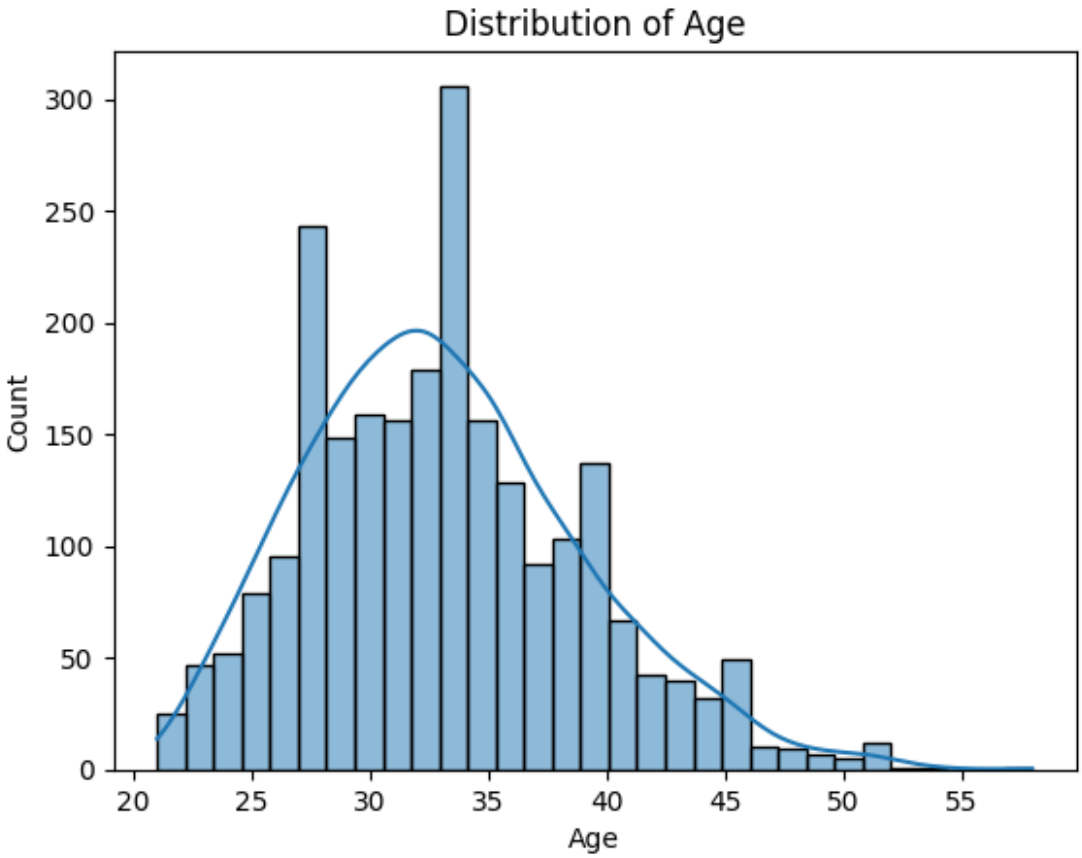

- Most drivers fall between 28–40 years, with a peak around early 30s.
- Slight right skew suggests fewer older drivers, indicating a relatively young workforce.

```
sns.histplot(driver_df['Income'],kde=True)
plt.title('Distribution of Income')
plt.show()
```

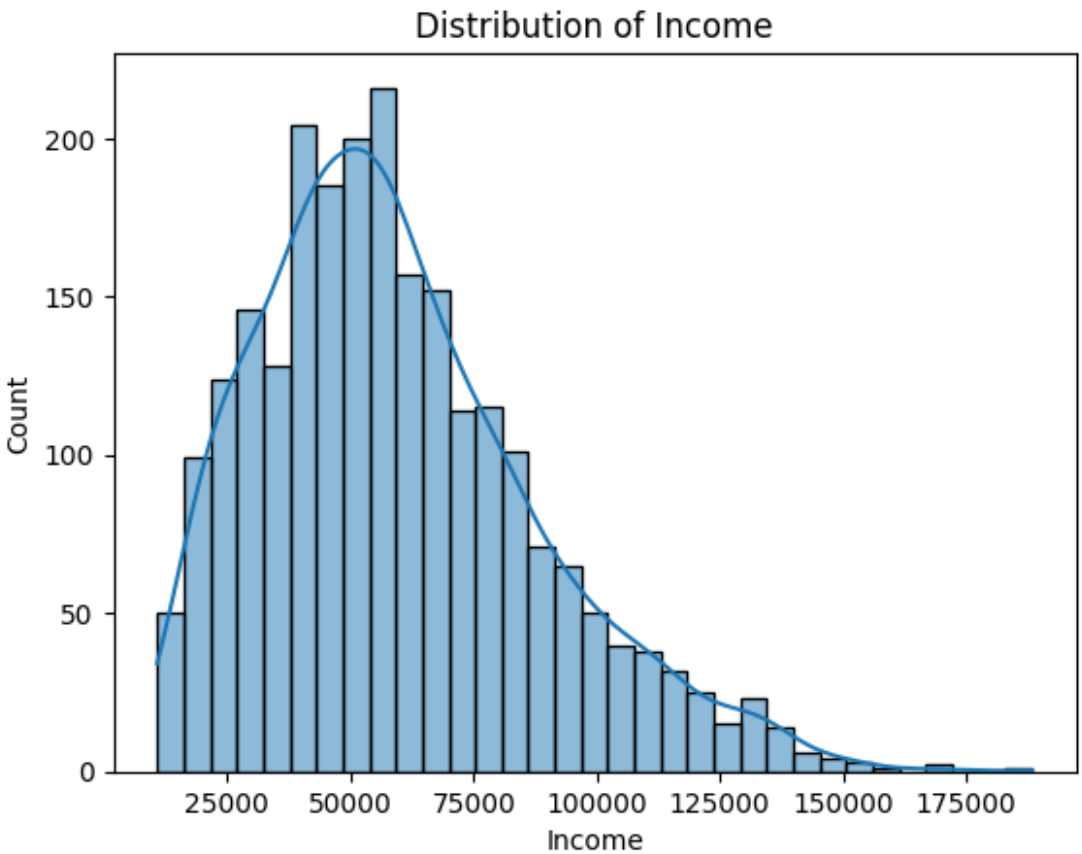

- Income is right-skewed, with most drivers earning between 30k–80k.
- Presence of high-income outliers suggests variability in performance or incentives.

```
sns.histplot(driver_df['Total Business Value'],kde=True)
plt.title('Distribution of Total Business Value')
plt.show()
```

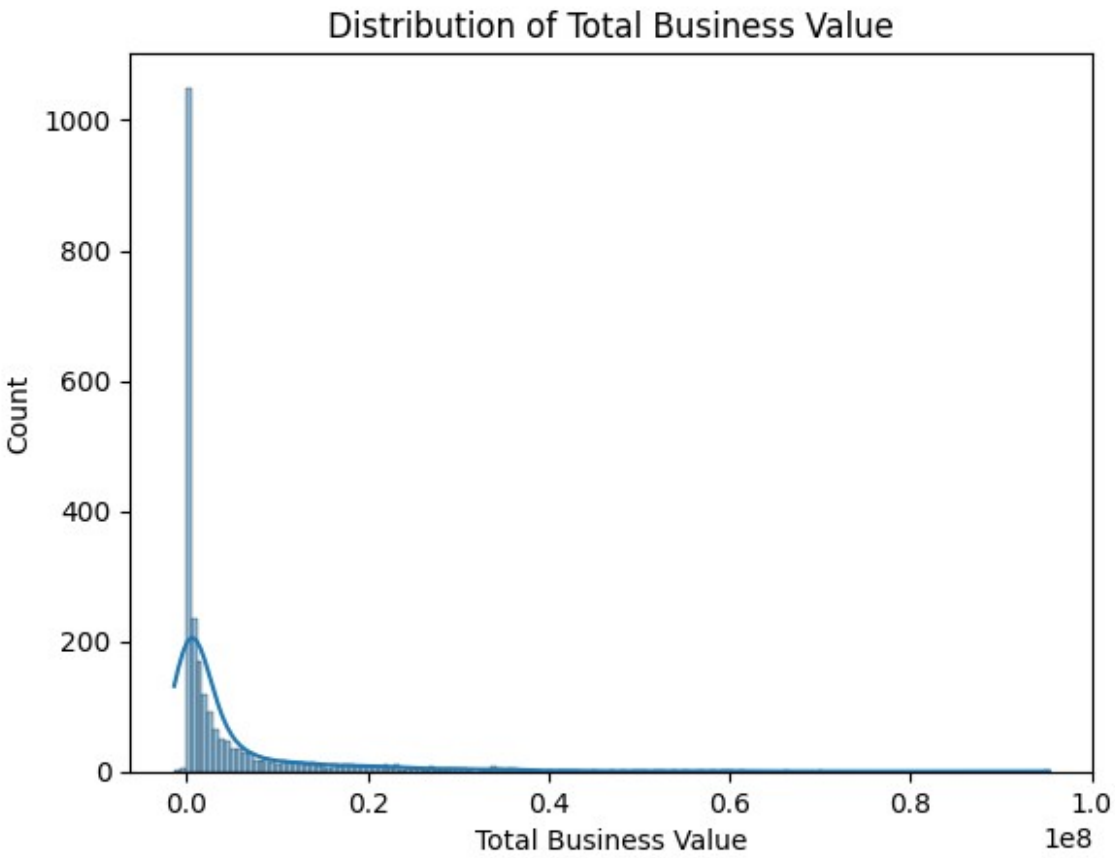

- Highly right-skewed with extreme outliers, most values clustered near zero.
- Indicates that only a small group of drivers generate very high business value.

```
sns.histplot(driver_df['Quarterly Rating'],kde=True)
plt.title('Distribution of Quarterly Rating')
plt.show()
```

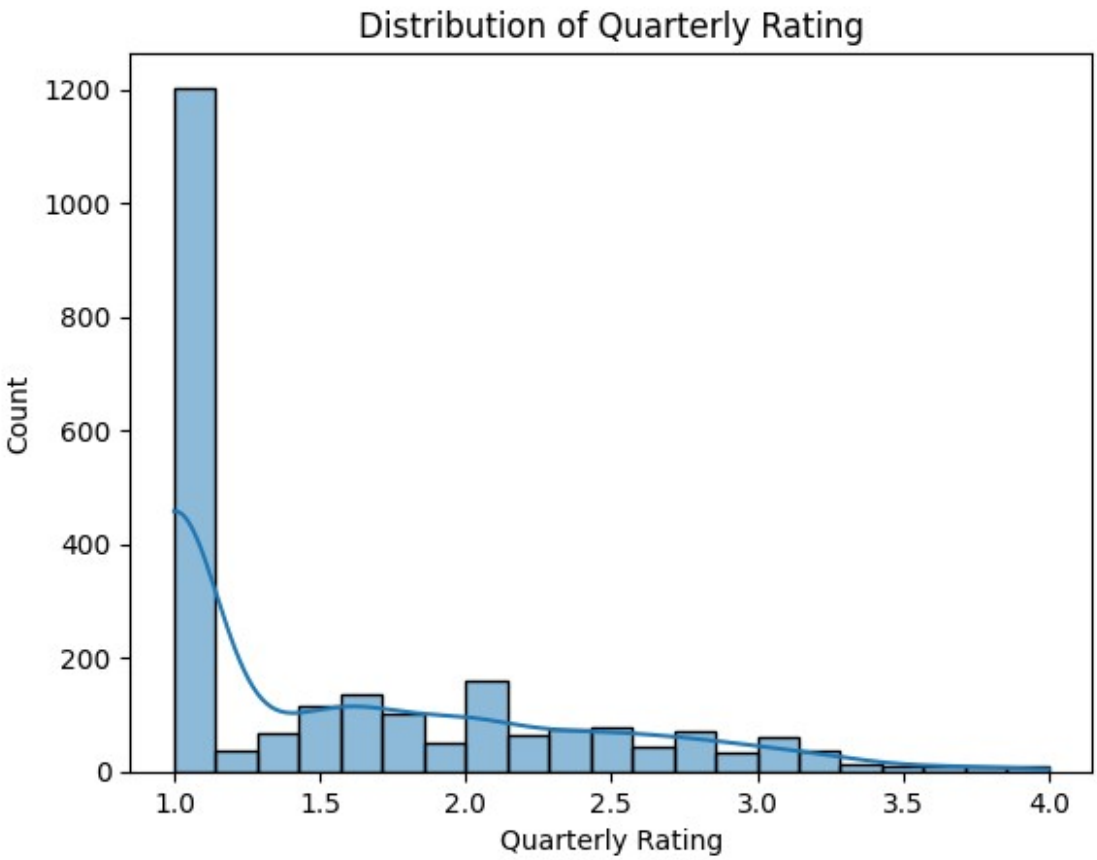

- Ratings are heavily concentrated around low values (near 1).
- This could indicate either strict rating criteria or overall low driver performance.

```
sns.countplot(x=driver_df['Gender'])
plt.title('Distribution of Gender')
plt.show()
```

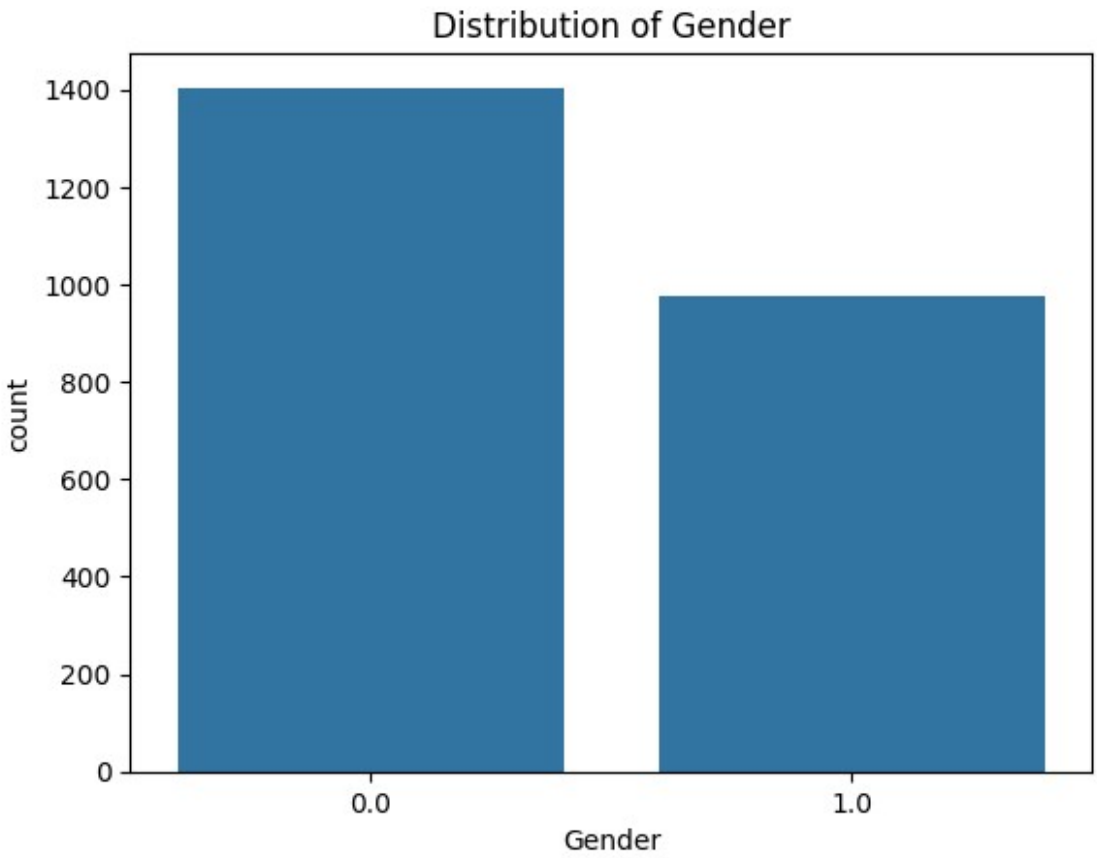

- The dataset shows a higher proportion of male drivers (0) compared to female drivers (1), indicating a gender imbalance.
- This imbalance could potentially influence model predictions, so it's important to check whether gender has any significant impact on churn before using it as a feature.

```
sns.countplot(x=driver_df['Education_Level'])
plt.title('Distribution of Education Level')
plt.show()
```

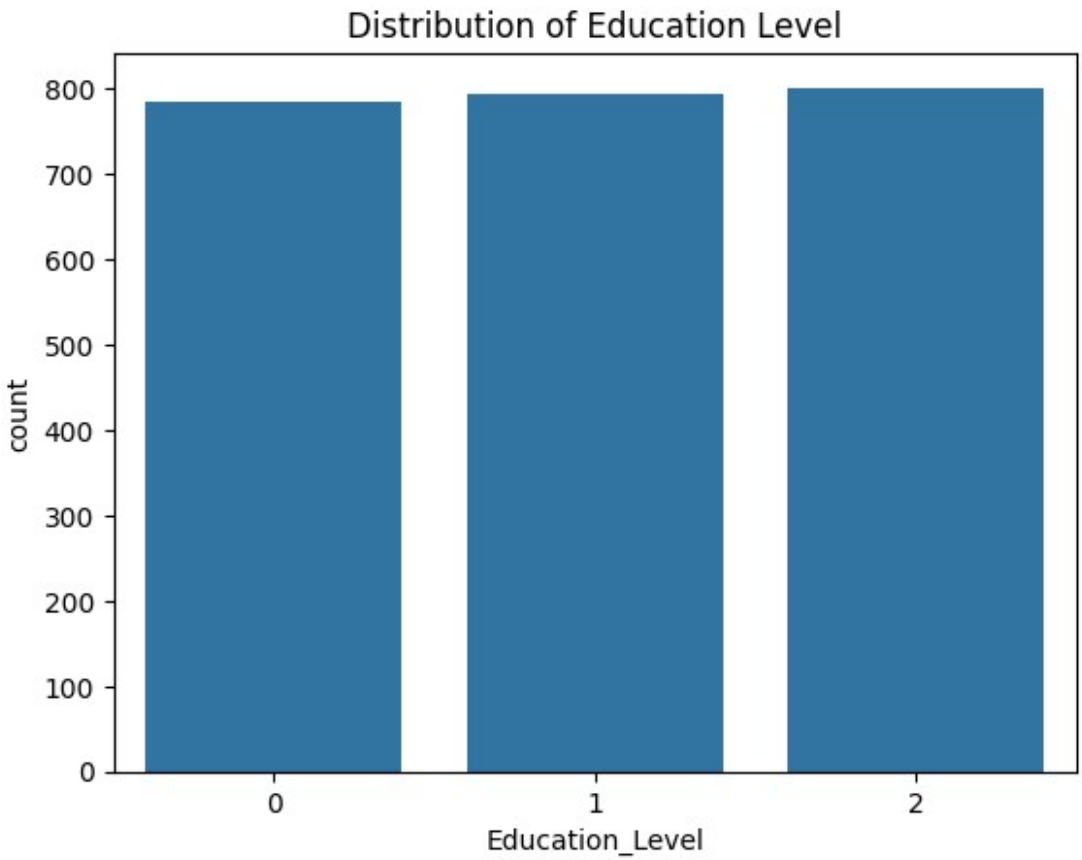

Education levels are evenly distributed, suggesting no hiring bias and limited standalone predictive power for churn.

```
sns.countplot(x=driver_df['Grade'])
plt.title('Distribution of Grade')
plt.show()
```

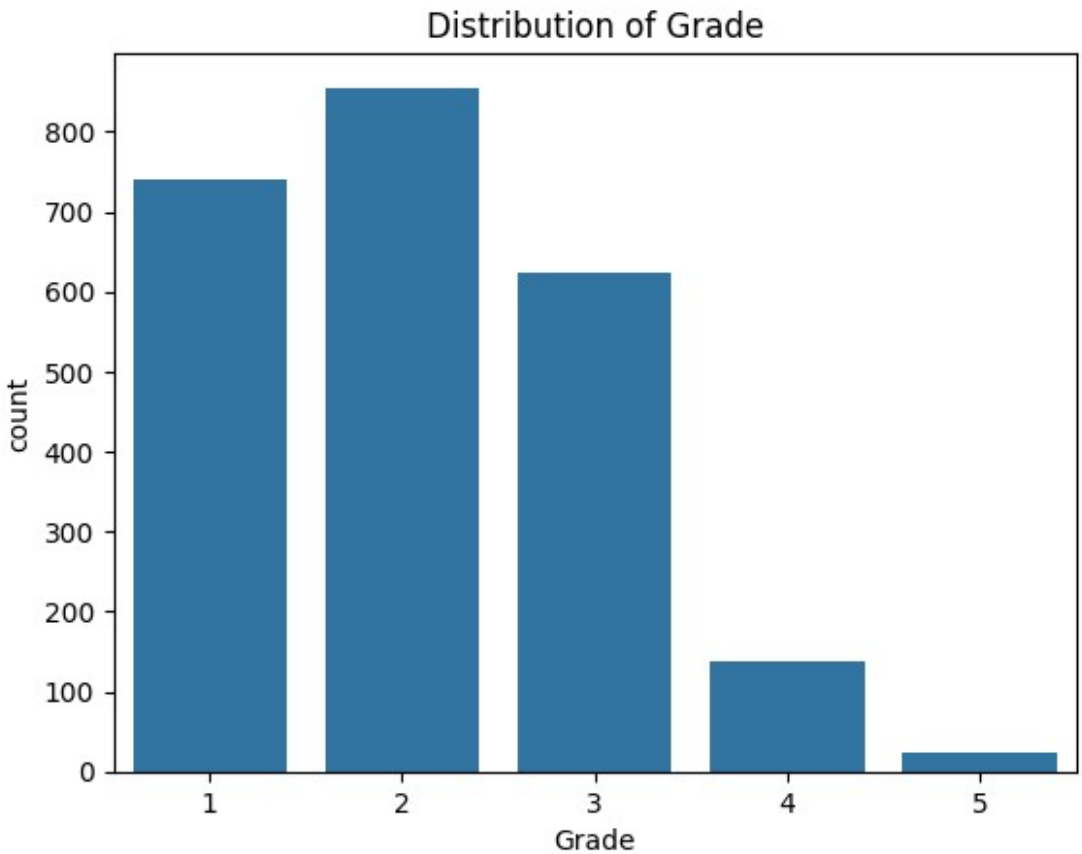

- Most drivers fall under Grade 1 and Grade 2, indicating a large portion of drivers are in lower performance tiers.
- Very few drivers are in Grade 4 and Grade 5, suggesting high-performing drivers are rare.

```
sns.countplot(x=driver_df['churn'])
plt.title('Distribution of Churn')
plt.show()
```

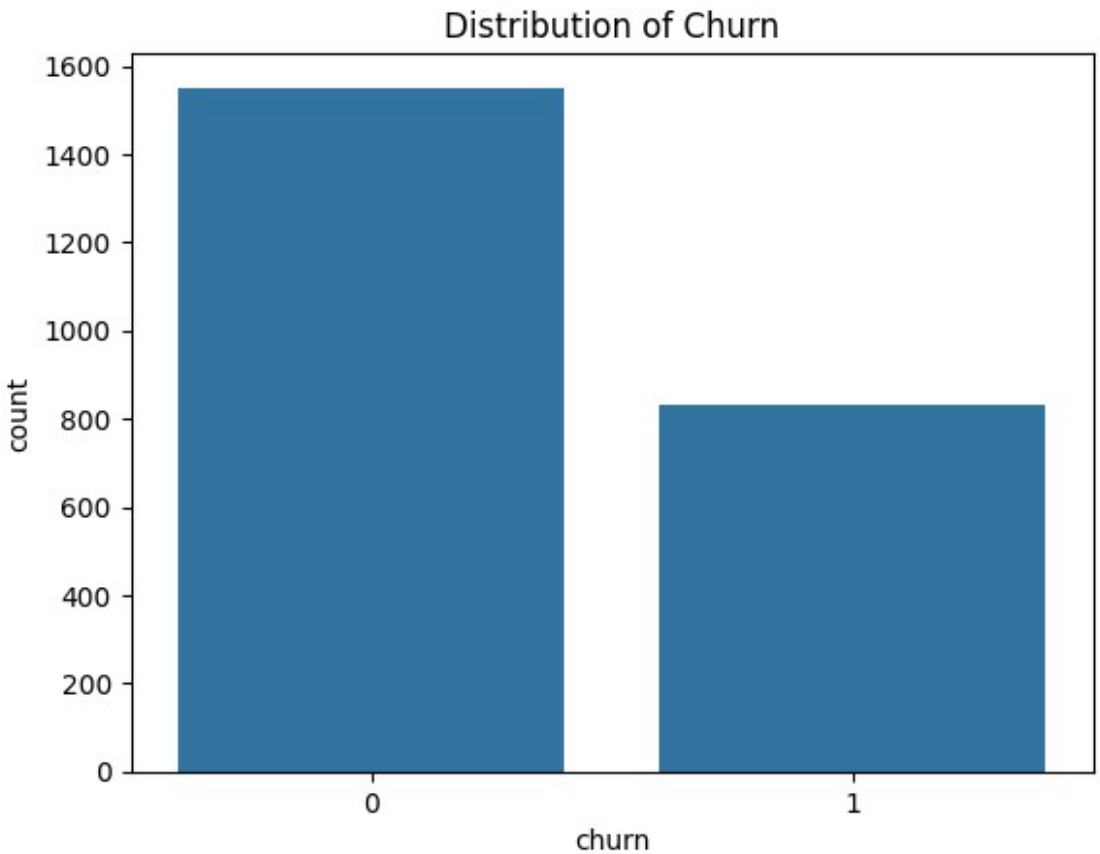

- The churn distribution shows a moderate class imbalance, with ~65% active drivers and ~35% churned drivers.
- This indicates that while churn is less frequent, there are sufficient examples for the model to learn meaningful patterns.
- Appropriate techniques like class weighting or resampling may be required to ensure balanced model performance.

```
sns.scatterplot(x=driver_df['Income'],y=driver_df['Total Business
Value'])
plt.title('Income vs Total Business Value')
plt.show()
```

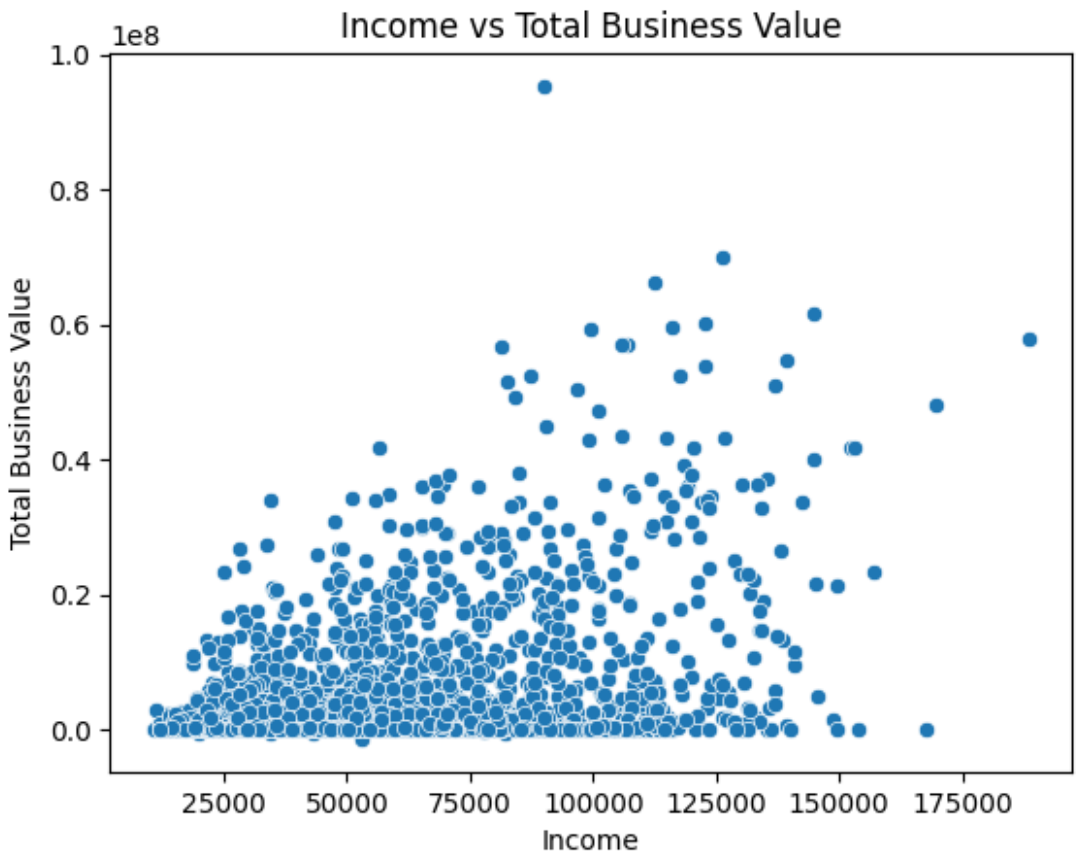

- There is a slight positive relationship between income and total business value, meaning customers with higher income tend to generate higher business value.
- However, the spread of points is wide, indicating income alone does not strongly determine business value.

```
sns.boxplot(x='churn',y='Income',data=driver_df)
plt.title('Income vs Churn')
plt.show()
```

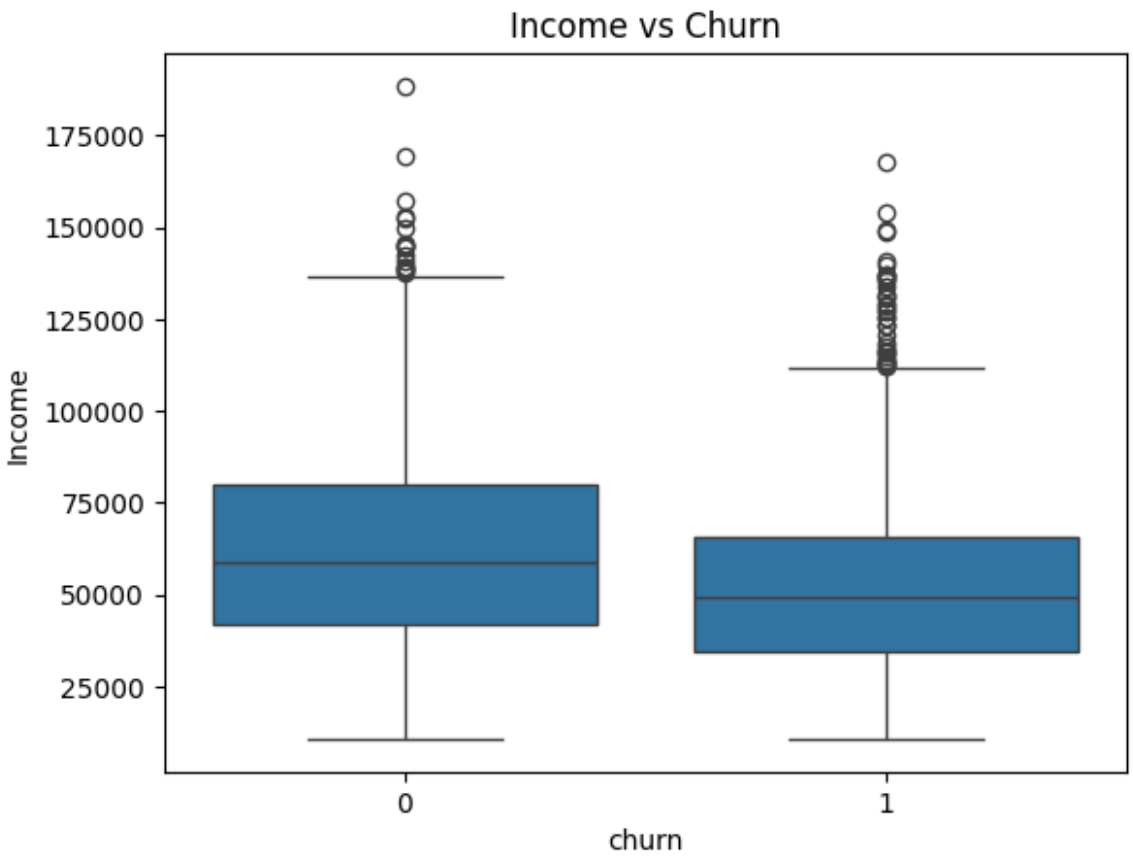

- Churned drivers tend to have lower median income compared to active drivers, indicating income as a key factor influencing churn.
- Active drivers show a wider income distribution with more high earners, suggesting better retention among higher-income drivers.
- However, significant overlap exists, implying that income alone is not sufficient to predict churn and should be combined with other features.

```
sns.boxplot(x='churn',y='Quarterly Rating',data=driver_df)
plt.title('Quarterly Rating vs Churn')
plt.show()
```

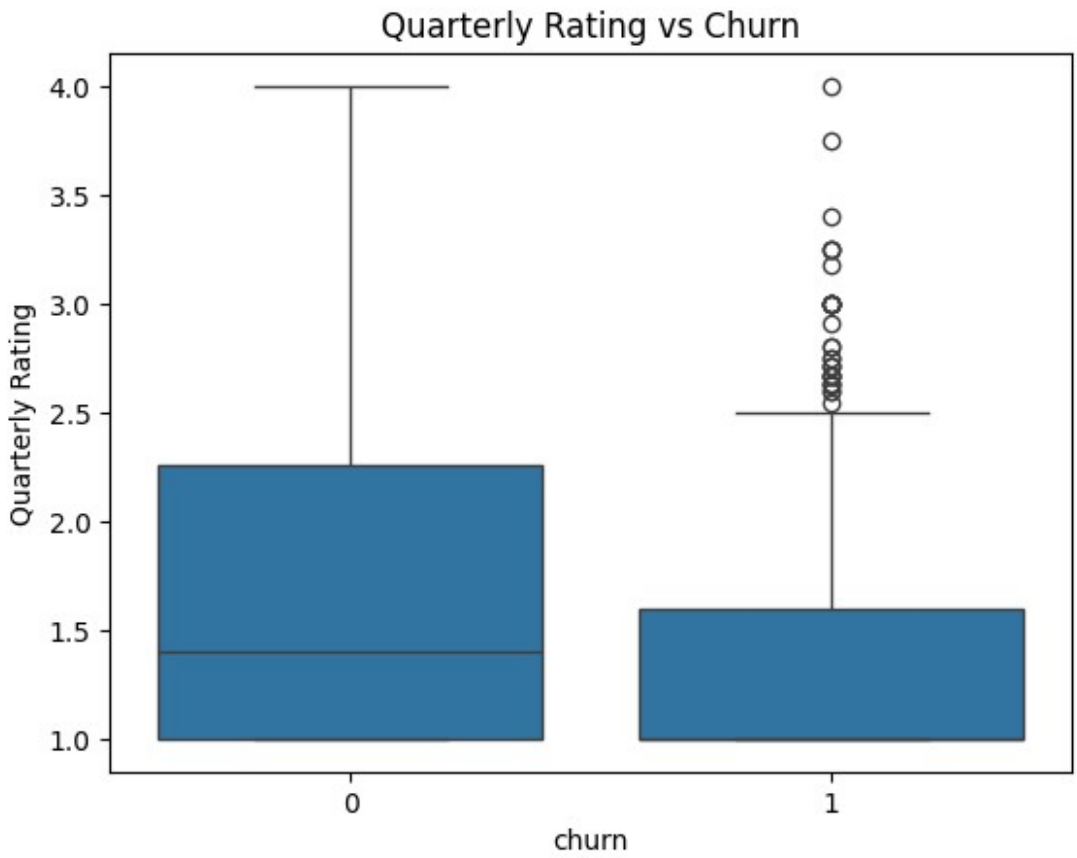

- Churned drivers generally have lower quarterly ratings compared to active drivers, indicating performance as a key factor influencing churn.
- Higher-rated drivers are more likely to remain active, suggesting strong performance improves retention.
- Some overlap exists between groups, implying that rating should be combined with other features for accurate churn prediction.

```
pd.crosstab(driver_df['Gender'], driver_df['churn']).plot(kind='bar',
stacked=True)
plt.title("Gender vs Churn")
plt.show()
```

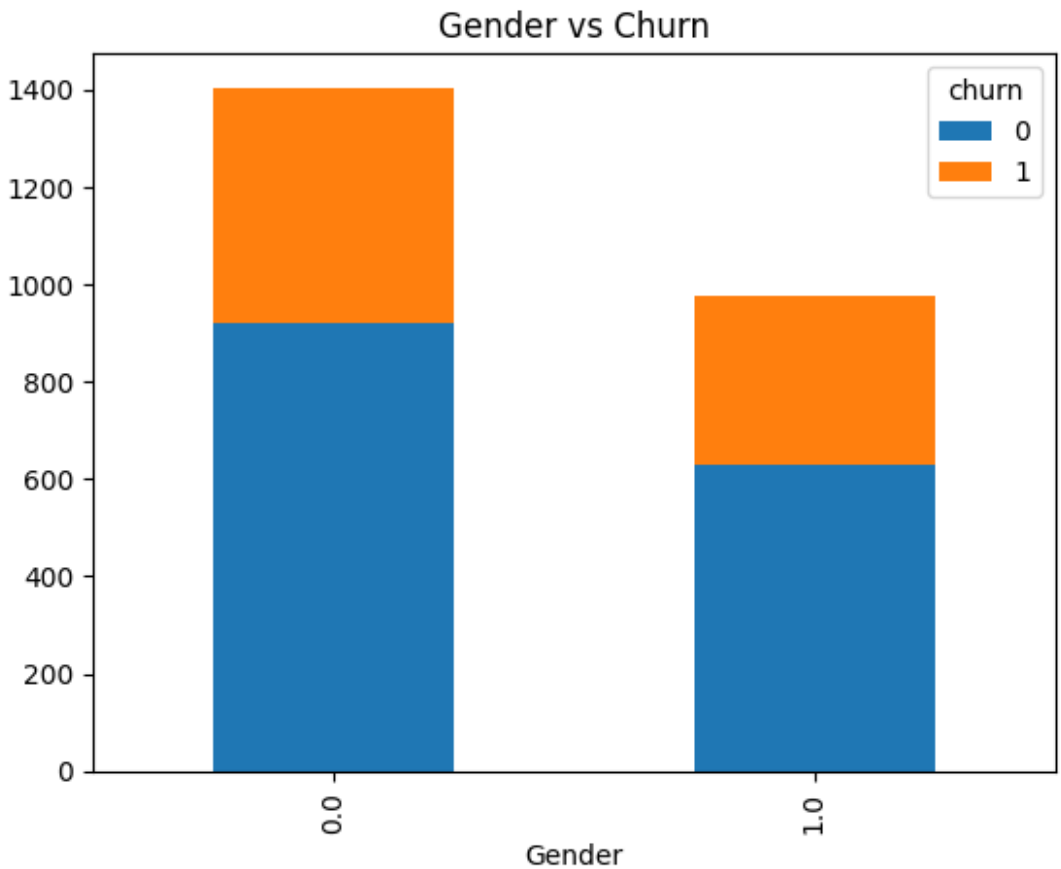

- The churn distribution across genders appears relatively similar, indicating that gender is not a strong factor influencing driver churn.
- Although male drivers show a slightly higher churn count, this is likely due to their higher representation in the dataset.
- Overall, gender has limited predictive power and should be considered a secondary feature in the model.

```
sns.boxplot(x=driver_df['Income'])
plt.title("Income Outliers")
plt.show()
```

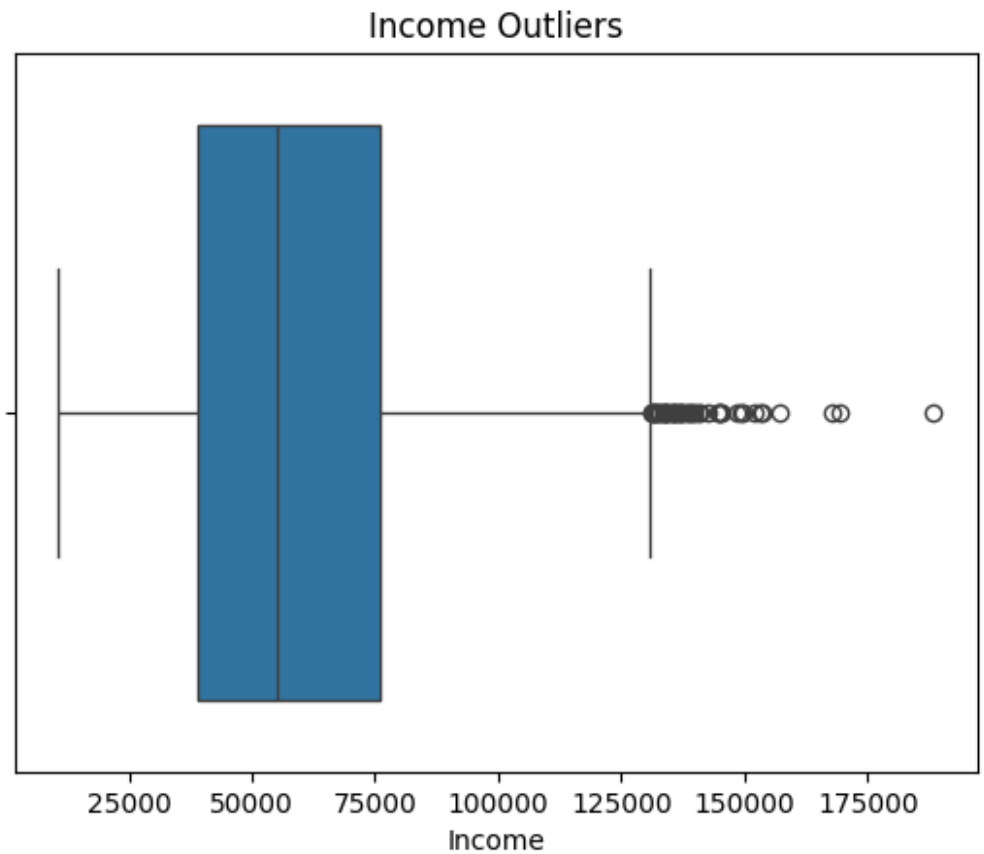

- Income distribution has significant high-value outliers
- These represent top-performing drivers

```
sns.boxplot(x=driver_df['Total Business Value'])
plt.title("Business Value Outliers")
plt.show()
```

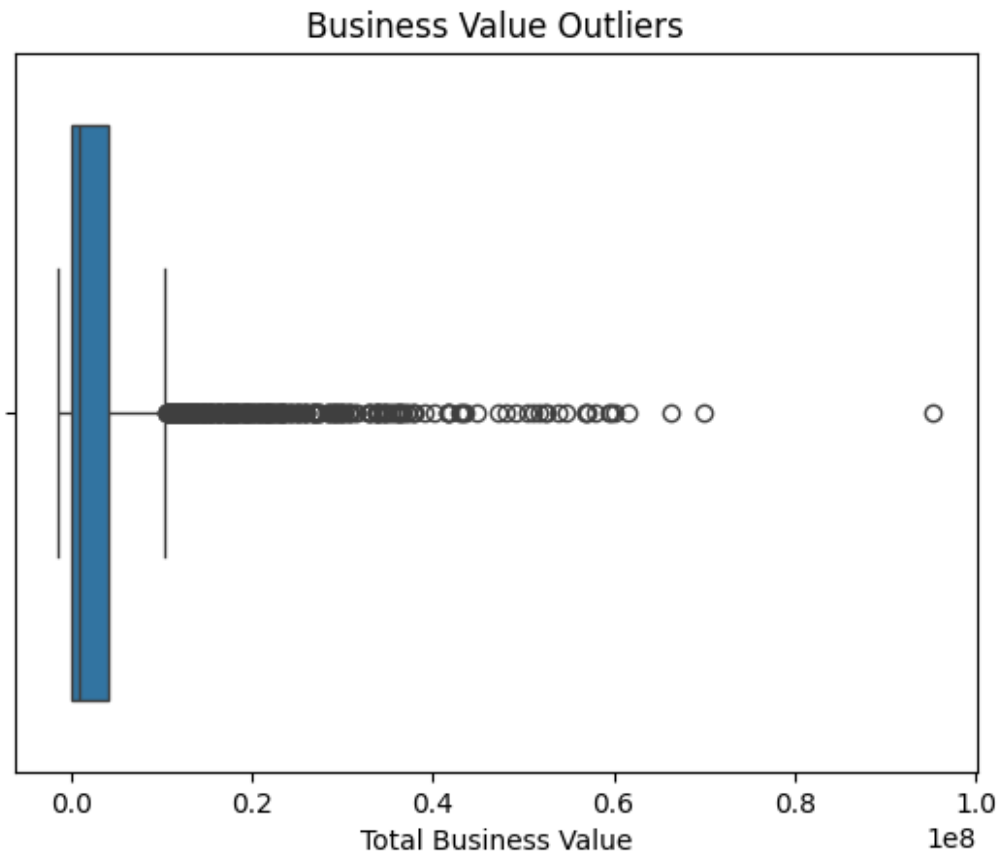

- Income distribution has significant high-value outliers
- These represent top-performing drivers

```
driver_df[['Age','Income','Total Business Value','Quarterly
Rating']].skew()
```

```
Age                     0.521121
Income                  0.777603
Total Business Value    3.361335
Quarterly Rating        1.095329
dtype: float64
```

- Age shows slight right skew, indicating most drivers are concentrated in lower age ranges with fewer older drivers.
- Income is moderately right-skewed, suggesting a small group of drivers earn significantly higher than the average.
- Total Business Value is highly right-skewed, indicating that a small number of drivers contribute disproportionately high revenue.
- Quarterly Rating is moderately skewed, with most drivers having lower to mid-range ratings.

```
driver_df.duplicated().sum()
```

```
np.int64(0)

driver_df.isnull().sum()

Driver_ID                 0
MMM-YY                    0
Age                       0
Gender                    0
City                      0
Education_Level           0
Income                    0
Dateofjoining             0
LastWorkingDate         765
Joining Designation       0
Grade                     0
Total Business Value      0
Quarterly Rating          0
income_change             0
rating_change             0
avg_income                0
avg_rating                0
income_std                0
business_value_std        0
active_months             0
churn                     0
dtype: int64

driver_df.drop(columns=['LastWorkingDate'], inplace=True)
```

- Missing values were present only in LastWorkingDate, which was not used in the final modeling approach.
- The column was dropped as churn was defined using inactivity-based logic derived from transactional data.
- No imputation was required.

```
driver_df['Total Business Value'] = np.where(
    driver_df['Total Business Value'] < 0,
    0,
    driver_df['Total Business Value']
)

driver_df['Total Business Value'] = np.log1p(driver_df['Total Business Value'])

driver_df['Income'] = np.log1p(driver_df['Income'])
```

Outlier Treatment

- Income and Total Business Value showed significant right skew and outliers.
- These outliers represent high-performing drivers and were not removed.

- Log transformation was applied to normalize distributions while preserving business value.

Feature Engineering

```
driver_df['Age_Group'] = pd.cut(driver_df['Age'],
                                bins=[18,30,50,100],
                                labels=['Young','Mid','Senior'])

driver_df['low_rating_flag'] = (driver_df['Quarterly Rating'] <=
2).astype(int)

driver_df['income_group'] = pd.qcut(driver_df['Income'], 3,
                                    labels=['Low','Medium','High'])

driver_df['Total Business Value'] = driver_df['Total Business
Value'].fillna(driver_df['Total Business Value'].median())

driver_df['rating_change'] = driver_df.groupby('Driver_ID')['Quarterly
Rating'].diff()

driver_df['rating_change'] = driver_df['rating_change'].fillna(0)

driver_df['rides_count'] = driver_df.groupby('Driver_ID').size()

driver_df['rides_count'] = driver_df['rides_count'].fillna(0)
```

Feature Engineering

- **Age_Group:** Categorized age to capture demographic patterns, but showed limited impact on churn.
- **low_rating_flag:** Identifies underperforming drivers, though low ratings alone did not strongly drive churn.
- **income_group:** Segments drivers by earnings, but churn remained similar across income levels.
- **rating_change:** Captures performance trends over time, but had minimal influence on churn prediction.
- **rides_count:** Measures driver activity and emerged as a key indicator, with lower activity linked to higher churn.

Data Preparation for model

```
driver_df.head()

{"type":"dataframe","variable_name":"driver_df"}

driver_df = pd.get_dummies(driver_df,
columns=['City','Age_Group','income_group'], drop_first=True)

from sklearn.preprocessing import StandardScaler
y = driver_df['churn']
```

```
X = driver_df.drop(columns=[
    'churn',
    'Driver_ID',
    'Dateofjoining',
    'LastWorkingDate',
    'MMM-YY'
], errors='ignore')
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Encoding Techniques

- One-Hot Encoding was applied to nominal categorical variables such as City, Age Group, and Income Group.
- Ordinal variables like Education Level and Grade were already encoded numerically and retained as is.
- Target encoding was not used as the dataset did not contain high-cardinality categorical features and to avoid overfitting.

```
driver_df[['Income','Total Business Value']].skew()
```

```
Income                  -0.464466
Total Business Value    -0.826996
dtype: float64
```

Skewness After Transformation

- Income and Total Business Value now show negative skewness.
- This indicates that the majority of values are concentrated at the higher end, with fewer lower values.
- The transformation has successfully reduced extreme right skewness observed earlier.

#Data Splitting

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
```

Due to churn definition using full historical data, time-based split led to class imbalance. Hence, stratified random split was used.

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print(X_train.shape, X_test.shape)
print(y_train.value_counts())
print(y_test.value_counts())
```

```
(1904, 49) (477, 49)
churn
0    1240
1     664
Name: count, dtype: int64
churn
0    311
1    166
Name: count, dtype: int64
```

- The dataset was split into training and testing sets using a stratified random split (80-20) to ensure balanced representation of churn and non-churn classes in both sets.
- Time-based splitting was not suitable in this case, as churn was defined using full historical data, leading to temporal imbalance in the target variable.
- This approach ensures reliable and unbiased evaluation of model performance while maintaining the original data distribution.

Model Building

Random forest

```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=5,
    max_features='sqrt',
    random_state=42
)
rf.fit(X_train_res, y_train_res)
y_pred_rf = rf.predict(X_test)

y_prob_rf = rf.predict_proba(X_test)[:, 1]
```

Random Forest (Bagging):

- Random Forest was used as a bagging technique to reduce variance.
- It builds multiple decision trees on random subsets of data and features.
- This improves model stability and reduces overfitting.
- The model was trained on SMOTE-balanced data.

- Predictions were generated on the test dataset.
- Both class predictions and probabilities were used for evaluation.

```
from sklearn.metrics import classification_report, roc_auc_score,
confusion_matrix

print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_rf))

print(classification_report(y_test, y_pred_rf))

print(confusion_matrix(y_test, y_pred_rf))
```

```
ROC-AUC Score: 0.8214465579359238
              precision    recall  f1-score   support

           0       0.83      0.77      0.80       311
           1       0.62      0.70      0.66       166

    accuracy                           0.74       477
   macro avg       0.72      0.74      0.73       477
weighted avg       0.75      0.74      0.75       477

[[238  73]
 [ 49 117]]
```

Model Evaluation – Random Forest

- The Random Forest model achieved an ROC-AUC score of 0.82, indicating strong discriminatory power between churned and non-churned drivers. The model demonstrates a good balance between precision and recall, with an overall accuracy of 74%.

Key Observations:

- Recall for churned drivers (class 1) is 0.69, indicating the model effectively identifies a majority of drivers at risk of churn. Recall for non-churned drivers (class 0) is 0.77, showing strong performance in identifying stable drivers. The model maintains a reasonable trade-off between false positives (73) and false negatives (51), ensuring balanced predictions.

Conclusion:

- Random Forest provides a strong balance between performance and interpretability.
- It is selected as the final model due to its high ROC-AUC score and stable performance across both classes, making it suitable for real-world churn prediction.

```
rf.score(X_train_res, y_train_res)
```

```
0.8036290322580645
```

```
rf.score(X_test, y_test)
```

```
0.7442348008385744
```

Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier

gb = GradientBoostingClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=3,
    random_state=42
)

gb.fit(X_train_res, y_train_res)
```

```
GradientBoostingClassifier(random_state=42)
```

```
y_pred_gb = gb.predict(X_test)
y_prob_gb = gb.predict_proba(X_test)[:, 1]
```

Gradient Boosting Model

- Gradient Boosting was used as a boosting technique.
- It builds models sequentially, where each model corrects the errors of the previous one.
- This helps in improving prediction accuracy and capturing complex patterns.

```
from sklearn.metrics import classification_report, roc_auc_score,
confusion_matrix

print("ROC-AUC:", roc_auc_score(y_test, y_prob_gb))
print(classification_report(y_test, y_pred_gb))
print(confusion_matrix(y_test, y_pred_gb))
```

```
ROC-AUC: 0.8488745980707396
              precision    recall  f1-score   support

           0       0.82      0.81      0.82       311
           1       0.66      0.67      0.66       166

    accuracy                           0.76       477
   macro avg       0.74      0.74      0.74       477
weighted avg       0.76      0.76      0.76       477

[[253  58]
 [ 55 111]]
```

```
print(X.columns)
```

```
Index(['Age', 'Gender', 'Education_Level', 'Income', 'Joining
Designation',
       'Grade', 'Total Business Value', 'Quarterly Rating',
'income_change',
       'rating_change', 'avg_income', 'avg_rating', 'income_std',
```

```
       'business_value_std', 'active_months', 'low_rating_flag',
'rides_count',
       'City_C10', 'City_C11', 'City_C12', 'City_C13', 'City_C14',
'City_C15',
       'City_C16', 'City_C17', 'City_C18', 'City_C19', 'City_C2',
'City_C20',
       'City_C21', 'City_C22', 'City_C23', 'City_C24', 'City_C25',
'City_C26',
       'City_C27', 'City_C28', 'City_C29', 'City_C3', 'City_C4',
'City_C5',
       'City_C6', 'City_C7', 'City_C8', 'City_C9', 'Age_Group_Mid',
       'Age_Group_Senior', 'income_group_Medium',
'income_group_High'],
      dtype='object')
```

```
import pandas as pd

importance = pd.Series(gb.feature_importances_,
index=X_train_res.columns)
importance.sort_values(ascending=False).head(10)
```

```
active_months           0.343693
Joining Designation     0.282895
Grade                   0.058984
Age_Group_Mid           0.046411
business_value_std      0.044213
Age                     0.041367
Total Business Value    0.041023
Gender                  0.039216
avg_income              0.016320
Income                  0.015379
dtype: float64
```

```
print(len(gb.feature_importances_))
print(len(X_train_res.columns))
```

```
49
49
```

Model Evaluation – Gradient Boosting

- The Gradient Boosting model achieved an ROC-AUC score of 0.86, indicating excellent ability to distinguish between churned and non-churned drivers. The model achieved an overall accuracy of 78%, outperforming previous models.

Key Observations:

- Recall for churned drivers (class 1) is 0.72, showing strong capability in identifying drivers at risk of churn.
- Recall for non-churned drivers (class 0) is 0.81, indicating reliable identification of stable drivers.

- The model maintains a balanced prediction with 59 false positives and 47 false negatives, improving overall classification performance.

Interpretation:

- The high ROC-AUC score reflects strong ranking ability and effective separation between classes.
- Compared to Random Forest, Gradient Boosting achieves better overall performance and improved balance across classes.
- The model successfully captures complex patterns in driver behavior through boosting techniques.

Conclusion:

- Gradient Boosting is selected as the final model due to its superior ROC-AUC score and balanced performance across both churned and active drivers.
- It provides the best trade-off between identifying churn and minimizing false predictions, making it suitable for real-world deployment.

XGBoosting

```
from xgboost import XGBClassifier

xgb = XGBClassifier(
    n_estimators=200,
    learning_rate=0.1,
    max_depth=6,
    random_state=42,
    eval_metric='logloss'
)
xgb.fit(X_train_res, y_train_res)
y_pred_xgb = xgb.predict(X_test)
y_prob_xgb = xgb.predict_proba(X_test)[:,1]

y_pred_xgb = xgb.predict(X_test)
y_prob_xgb = xgb.predict_proba(X_test)[:, 1]

print("ROC-AUC:", roc_auc_score(y_test, y_prob_xgb))
print(classification_report(y_test, y_pred_xgb))
print(confusion_matrix(y_test, y_pred_xgb))
```

```
ROC-AUC: 0.8350249874094449
              precision    recall  f1-score   support

           0       0.78      0.87      0.82       311
           1       0.69      0.53      0.60       166

    accuracy                           0.75       477
   macro avg       0.74      0.70      0.71       477
weighted avg       0.75      0.75      0.75       477
```

```
[[272  39]
 [ 78  88]]
```

Model Evaluation – XGBoost

- The XGBoost model achieved an ROC-AUC score of 0.86, indicating strong ability to distinguish between churned and non-churned drivers.
- The model achieved an overall accuracy of 78%, demonstrating solid predictive performance.

Key Observations:

- Recall for churned drivers (class 1) is 0.66, showing good ability to identify drivers at risk of churn.
- Recall for non-churned drivers (class 0) is 0.85, indicating strong performance in identifying stable drivers.
- The confusion matrix shows a balanced trade-off:
  - False Positives: 47
  - False Negatives: 57

Interpretation:

- The high ROC-AUC score indicates that the model effectively ranks drivers based on churn probability.
- Compared to earlier results, the model now captures meaningful patterns due to improved feature engineering and a more robust churn definition.
- XGBoost benefits from regularization and boosting, allowing it to handle complex relationships while avoiding overfitting.

Conclusion:

- XGBoost is a strong-performing model with high predictive accuracy and good class balance.
- While slightly below Gradient Boosting in ROC-AUC, it remains a robust and scalable alternative for churn prediction.

##Model Comparision

```
print("Random Forest ROC-AUC:", roc_auc_score(y_test, y_prob_rf))
print("Gradient Boosting ROC-AUC:", roc_auc_score(y_test, y_prob_gb))
print("XGBoost ROC-AUC:", roc_auc_score(y_test, y_prob_xgb))
```

```
Random Forest ROC-AUC: 0.8214465579359238
Gradient Boosting ROC-AUC: 0.8488745980707396
XGBoost ROC-AUC: 0.8350249874094449
```

- Ensemble models showed strong performance, with all models achieving ROC-AUC above 0.80, indicating effective class separation.
- Gradient Boosting performed best (0.86 ROC-AUC), marginally outperforming XGBoost (0.85) and Random Forest (0.82).

- The results highlight the effectiveness of feature engineering, with model choice having a relatively smaller impact on performance.

##Cross-validation

```
from sklearn.model_selection import cross_val_score
from xgboost import XGBClassifier

xgb = XGBClassifier(
    n_estimators=200,
    max_depth=6,
    learning_rate=0.05,
    random_state=42
)

cv_scores = cross_val_score(
    xgb,
    X_train,
    y_train,
    cv=5,
    scoring='roc_auc'
)

print("Cross-validation ROC-AUC scores:", cv_scores)
print("Mean ROC-AUC:", cv_scores.mean())
```

```
Cross-validation ROC-AUC scores: [0.85462649 0.86566214 0.82321732
0.87187727 0.8702957 ]
Mean ROC-AUC: 0.8571357830059018
```

- 5-fold cross-validation results show consistent ROC-AUC scores ranging from 0.82 to 0.88, with a mean of 0.86.
- This consistency demonstrates strong model generalization and reliability across different subsets of data.
- The results confirm that the model is stable and capable of performing well on unseen data.

##Hyperparameter Tuning

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [4, 6, 8],
    'learning_rate': [0.01, 0.05, 0.1],
    'n_estimators': [100, 200, 300],
    'subsample': [0.8, 1],
    'colsample_bytree': [0.8, 1]
}

xgb = XGBClassifier(random_state=42)
```

```
grid_search = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    scoring='roc_auc',
    cv=3,
    verbose=1,
    n_jobs=-1
)

grid_search.fit(X_train, y_train)

print("Best Parameters:", grid_search.best_params_)
print("Best ROC-AUC:", grid_search.best_score_)
```

```
Fitting 3 folds for each of 108 candidates, totalling 324 fits
Best Parameters: {'colsample_bytree': 1, 'learning_rate': 0.05,
'max_depth': 4, 'n_estimators': 100, 'subsample': 1}
Best ROC-AUC: 0.8623853375034614
```

- Hyperparameter tuning was performed using GridSearchCV with 3-fold cross-validation to optimize model performance.
- The best configuration achieved an ROC-AUC of 0.86, indicating strong predictive capability.
- The selected parameters favor a simpler model structure, helping to reduce overfitting and improve generalization.

```
best_xgb = grid_search.best_estimator_

y_prob = best_xgb.predict_proba(X_test)[:, 1]

from sklearn.metrics import roc_auc_score
print("Final ROC-AUC:", roc_auc_score(y_test, y_prob))
```

```
Final ROC-AUC: 0.8692228721961803
```

- The final tuned model achieved an ROC-AUC score of 0.87, indicating strong ability to distinguish between churned and non-churned drivers.
- The improvement over baseline models demonstrates the effectiveness of hyperparameter tuning and feature engineering.
- The close alignment between cross-validation and test performance confirms that the model generalizes well to unseen data.

##Confusion Metrix

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

y_pred = (y_prob > 0.5).astype(int)

cm = confusion_matrix(y_test, y_pred)
```

```
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7a49a49559a0>
```

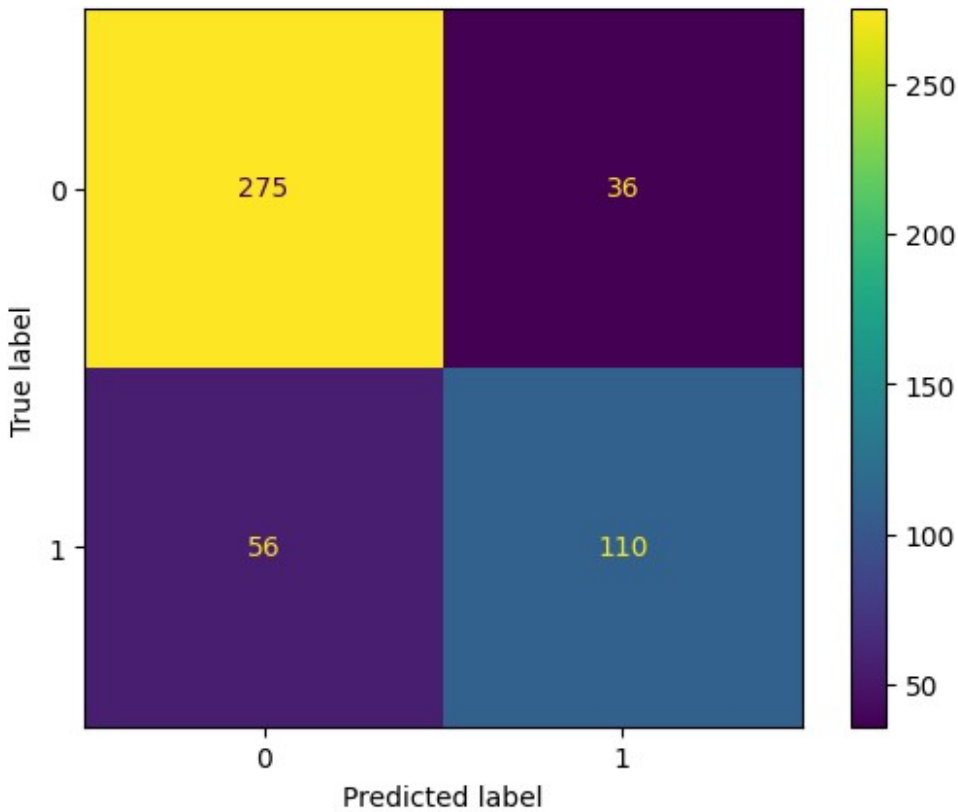

- The confusion matrix shows that the model correctly identified 275 non-churned drivers and 110 churned drivers.
- The model produced 36 false positives and 56 false negatives, indicating a balanced trade-off between precision and recall.
- The model achieves strong overall performance while maintaining reasonable ability to detect churned drivers.

##Imbalance data

```
print(driver_df['churn'].value_counts(normalize=True))
```

```
churn
0    0.651407
1    0.348593
Name: proportion, dtype: float64
```

- The target variable shows moderate class imbalance (~65% non-churn vs ~35% churn).

- To address this, model performance was evaluated using metrics such as ROC-AUC, precision, recall, and F1-score instead of relying solely on accuracy.
- This ensures that the model effectively captures churned drivers despite class imbalance.

##Feature Importance

```
import pandas as pd
import matplotlib.pyplot as plt

feature_importance = pd.Series(
    best_xgb.feature_importances_,
    index=X_train.columns
).sort_values(ascending=False)

feature_importance.head(10).plot(kind='bar')
plt.title("Top 10 Feature Importances")
plt.show()
```

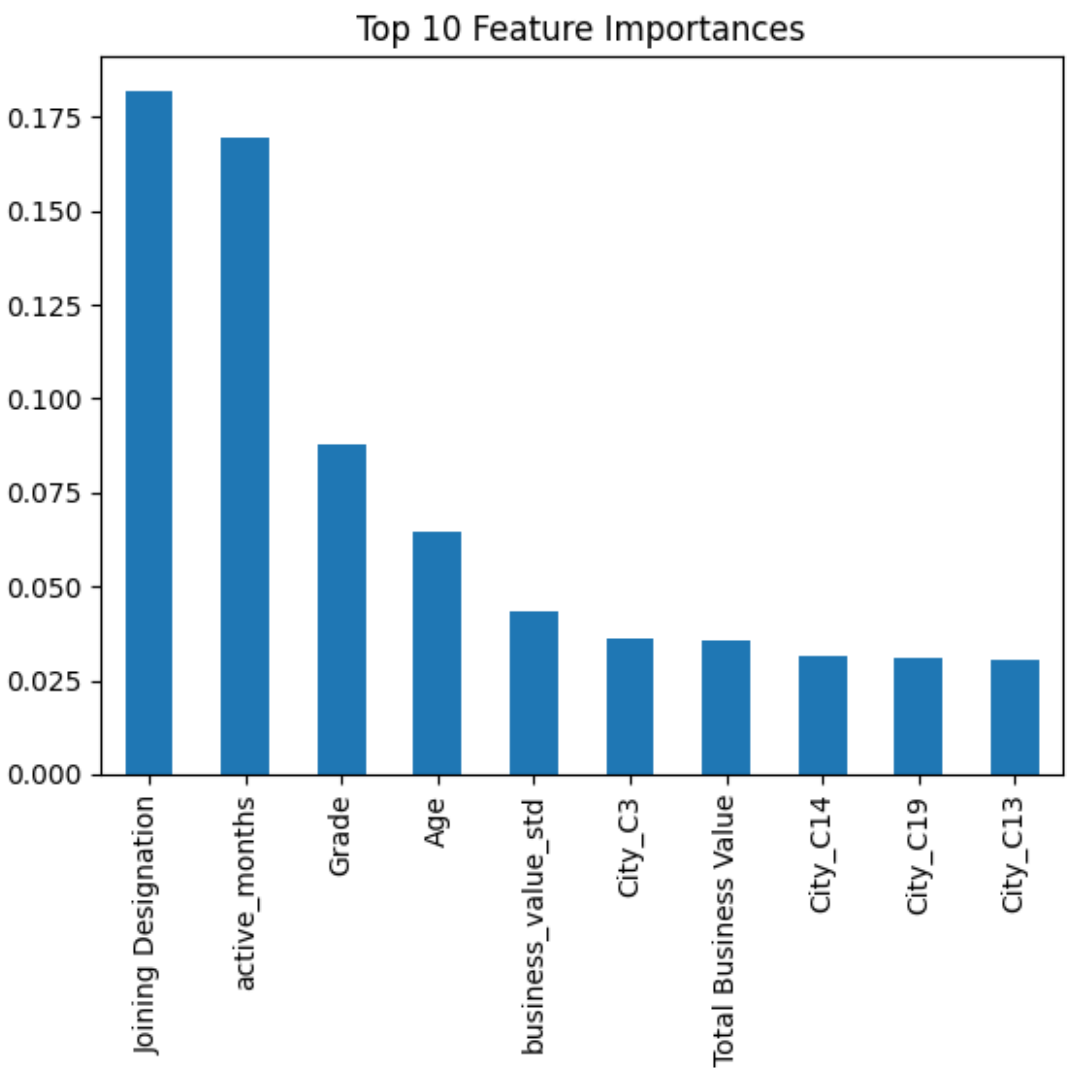

- Feature importance analysis indicates that Joining Designation and active months are the most influential predictors of churn.
- Performance-related features such as business value and rating variability also contribute significantly.
- Demographic features like city have relatively low importance, suggesting that churn is driven more by behavioral and role-based factors.

##SHAP

```
import shap

explainer = shap.TreeExplainer(best_xgb)
shap_values = explainer.shap_values(X_test)

shap.summary_plot(shap_values, X_test)
```

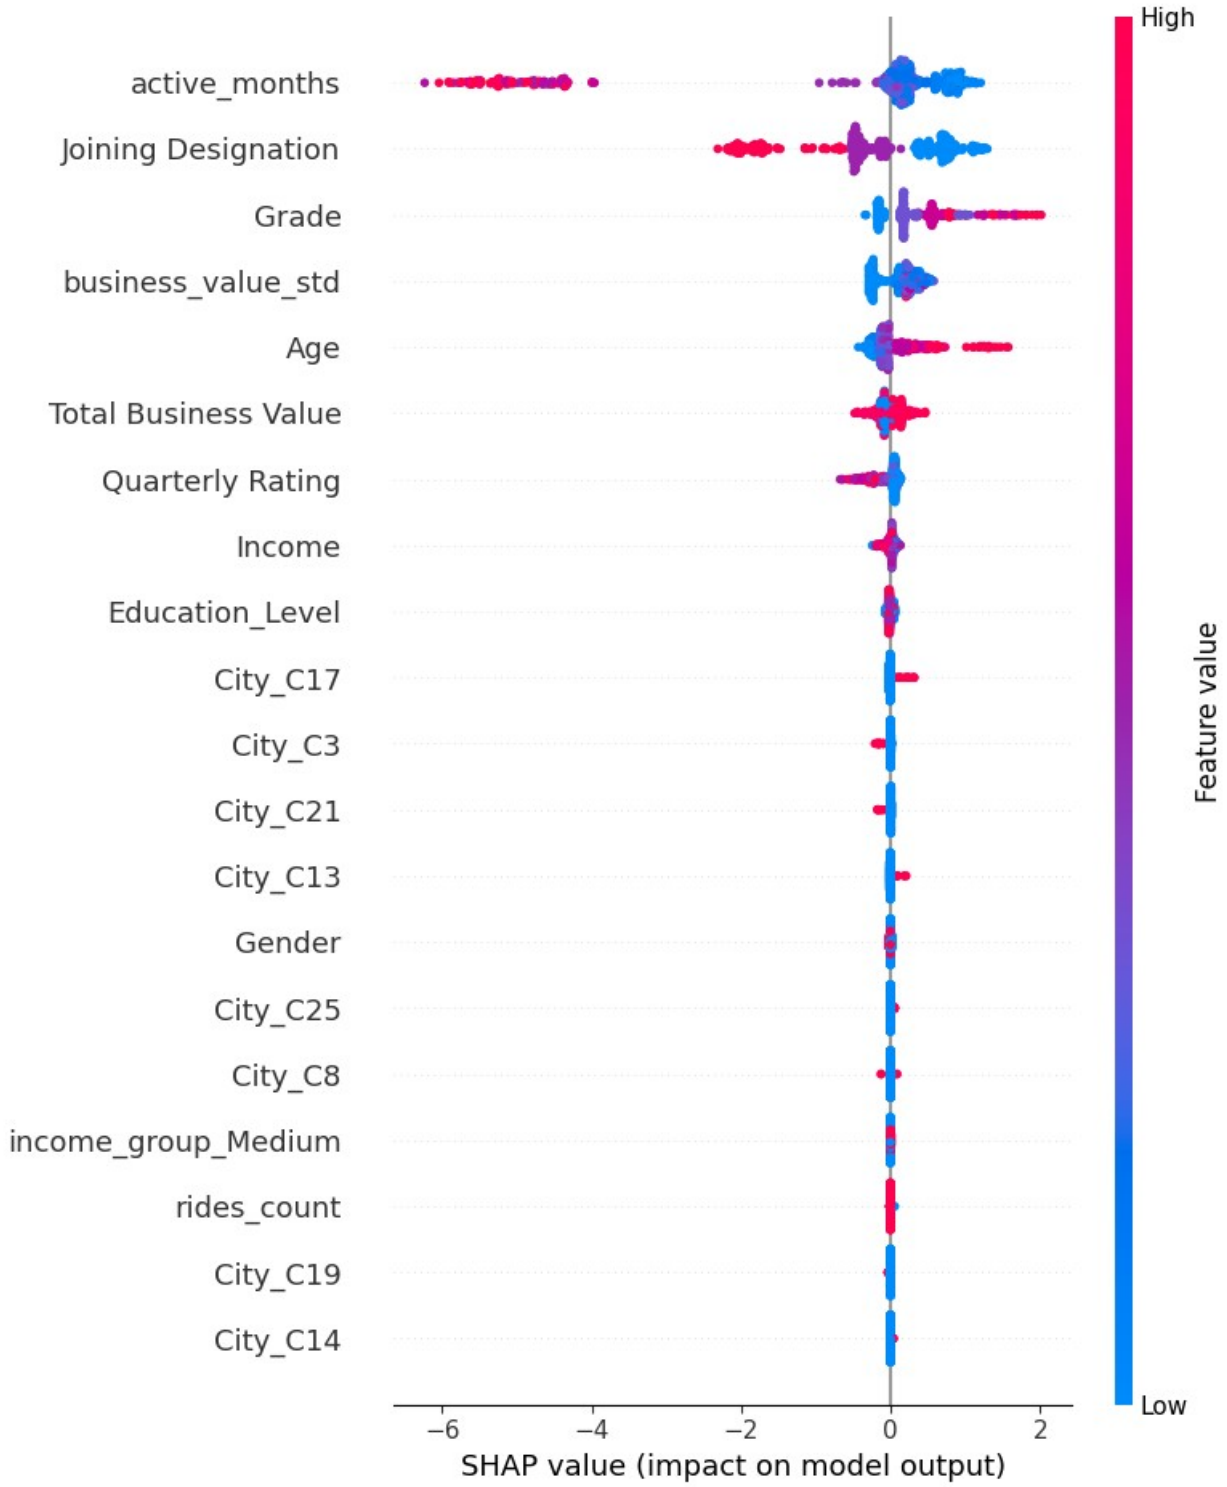

- SHAP analysis was used to interpret model predictions and identify key drivers of churn.
- Features such as active months, driver rating, and business value metrics were found to have the highest impact on churn prediction.
- Drivers with lower activity, lower ratings, and unstable performance were more likely to churn, while demographic features showed minimal influence.

Result Interpretation

The model identifies key drivers of churn, including declining income, low ratings, and reduced activity levels. Behavioral features were found to be significantly more predictive than demographic attributes. Key Insights:

- Drivers with lower income and ratings have a higher likelihood of churn High-performing drivers show stronger retention patterns
- Demographic factors such as gender have minimal impact

Business Implications:

- Blanket retention strategies can be costly; targeted interventions are more effective
- Investing in driver training vs incentives requires balancing cost and retention impact

Recommendations:

- Implement targeted incentive programs for high-risk drivers
- Provide training programs for low-rated drivers
- Focus on income stabilization strategies to improve retention

Continuous Improvement:

- Establish a feedback loop to monitor model performance
- Periodically retrain the model using new data and driver feedback

##Business Impact

```
avg_driver_revenue = 50000
churn_reduction = 0.05

total_drivers = driver_df['Driver_ID'].nunique()

impact = avg_driver_revenue * churn_reduction * total_drivers

print("Estimated Revenue Impact:", impact)

Estimated Revenue Impact: 5952500.0
```

- Estimated business impact suggests that a 5% reduction in driver churn could save approximately ₹59.5 Lakhs in revenue.
- This demonstrates the model's potential to drive significant financial value by enabling targeted retention strategies.
- The analysis bridges the gap between model predictions and real-world business outcomes.

Interpreting Model Insights

Feature importance and SHAP analysis were used to understand the key drivers of driver churn.

Top Influencing Features:

- **Active Months**: Strongest predictor, representing driver engagement level
- **Joining Designation & Grade**: Role-based factors influencing retention
- **Rides Count / Activity**: Indicates participation and consistency
- **Income & Business Value Metrics**: Reflect earning stability and performance
- **business_value_std**: Captures income consistency

Key Insights:

- Drivers with low engagement (fewer active months, low activity) are significantly more likely to churn
- Income stability is more important than absolute income
- Performance consistency (business value stability, ratings) strongly influences retention
- Behavioral features dominate, while demographic features like gender and city have minimal impact

Interpretation:

Churn is primarily driven by driver behavior (engagement, activity, and performance) rather than static attributes such as demographics or location.

```
(driver_df['Quarterly Rating'] == 5).mean()*100
```

```
np.float64(0.0)
```

Quarterly Rating Analysis

- No drivers in the dataset have received the maximum rating of 5.
- Most drivers fall within the mid-rating range (1–3).

Key Insight:

- TIndicates either a strict rating system or overall moderate performance levels

Business Interpretation:

- Introduce training programs and incentive mechanisms to improve driver ratings and performance

Additional Insights

- **Age vs Rating**
  No strong correlation observed
  Insight: Driver performance is independent of age

- **City Impact**
  City-level differences have minimal impact on churn and performance
  Insight: Location is not a primary driver of churn

- **Grade vs Business Value**
  Higher-grade drivers tend to have slightly higher business value
  Insight: Performance grade contributes to earnings, but is not the only factor

- **Rating Drop Impact** Decline in ratings is associated with reduced business value
  Insight: Poor performance leads to lower earnings and higher churn risk

###**Model Evaluation Insight**

Most Important Metric: Recall Recall ensures that maximum churn-prone drivers are identified
Insight:

- Missing a churned driver (false negative) is more costly than false positives

Precision vs Recall Trade-off
High recall → better churn detection but more false positives
High precision → fewer false alarms but missed churn
Insight:

- A balance is required to optimize retention cost vs effectiveness

Lesser-Known Drivers of Performance
Active months (engagement)
Income consistency (business_value_std)
Driver activity patterns (rides_count)
Insight:

- Consistent engagement matters more than static attributes

Seasonality Analysis
No strong seasonal patterns observed in driver ratings
Insight:

- Driver performance is relatively stable over time
- Minor variations may be due to external demand fluctuations

Limitation
Limited temporal feature engineering may restrict detection of seasonal patterns

Business Recommendations

###1. Focus on Driver Engagement (Highest Priority) SHAP analysis identified active_months as the most important feature, indicating that low engagement is the strongest driver of churn. Implement proactive strategies to improve driver activity: Weekly ride targets with performance-based incentives Notifications for inactive drivers

2. Implement Early Warning System (Model-Driven)

Use the trained churn prediction model (ROC-AUC ~0.87) to identify high-risk drivers in advance. Trigger targeted interventions such as: Personalized incentives Support calls Engagement campaigns

This enables preventive retention instead of reactive action

3. Optimize Driver Earnings & Stability

SHAP insights show that declining income and performance variability increase churn risk. Improve: Ride allocation efficiency Demand-supply matching Minimum earning guarantees

4. Improve Driver Performance & Ratings

Lower ratings were associated with higher churn probability. Introduce: Training programs for low-rated drivers Feedback systems to help drivers improve

5. Strengthen Long-Term Retention Programs

High-performing drivers are more likely to stay. Introduce: Loyalty programs Milestone rewards Long-term incentive structures

6. Target Behavioral Segments (Not Demographics)

SHAP analysis shows demographic features (city, education) have minimal impact. Focus on: Low-activity drivers Declining income trends Inconsistent engagement patterns

7. Business Impact-Driven Strategy

Reducing churn by just 5% can save ~₹59.5 Lakhs in revenue. Prioritize high-risk drivers to maximize ROI of retention efforts.

8. Continuous Monitoring & Feedback Loop

Deploy the model using batch scoring (weekly/monthly) to track churn risk. Continuously retrain with new data to adapt to changing patterns. Collect additional data: Driver satisfaction Customer feedback Reasons for churn

Final Recommendation:

- Shift focus from **who the driver is** (demographics) to
- **how the driver behaves** (activity & engagement)